

Forecasting Google Trends with VAR and ARIMA Models

Pasquale Gravante

2024-06-24

INTROUDCTION

This project aims to forecast Google Trends data for search terms using Vector Autoregression (VAR) and Autoregressive Integrated Moving Average (ARIMA) models. The analysis will involve data collection, preprocessing, stationarity testing, model selection, diagnostics, and forecasting. By the end of this project, we will compare the performance of VAR and ARIMA models in terms of their forecasting accuracy.

For the first part of the project, the data of the last 5 years for the “AI” search term is going to be used and then also the data for the “Chat GPT” search term is going to be used.

FIRST TIME SERIES

DATA COLLECTION AND PRE-PROCESSING

```
library(gtrendsR)
library(vars)
library(forecast)
library(ggplot2)
library(urca)
library(xts)
library(readr)
library(tidyverse)
library(lmtest)
library(tseries)
library(fUnitRoots)
```

```
library(quantmod)
library(kableExtra)
library(lubridate)
library(formattable)
```

The necessary libraries are loaded, and the data is imported from a CSV file. The first few rows of the dataset are displayed to ensure it has been loaded correctly.

```
# Load the CSV file
trend1_data <- read_csv("AITrend.csv")

# Check the data
head(trend1_data)
```

```
# A tibble: 6 x 2
  Week      AI
  <date>   <dbl>
1 2019-05-26 14
2 2019-06-02 14
3 2019-06-09 14
4 2019-06-16 14
5 2019-06-23 14
6 2019-06-30 14
```

PRE-PROCESSING

The 'Week' column is converted to Date format for time series analysis. The structure of the dataset is checked to confirm the conversion.

```
# Convert the 'week' column to Date type
trend1_data$Week <- as.Date(trend1_data$Week, format = "%Y-%m-%d")

# Check the data types to ensure the date conversion was successful
str(trend1_data)
```

```
spc_tbl_ [262 x 2] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
 $ Week: Date[1:262], format: "2019-05-26" "2019-06-02" ...
 $ AI  : num [1:262] 14 14 14 14 14 14 13 13 13 13 ...
 - attr(*, "spec")=
```

```

.. cols(
..   Week = col_date(format = ""),
..   AI = col_double()
.. )
- attr(*, "problems")=<externalptr>

```

The dataset is then converted into an xts object, which is a common format for time series data since it facilitates time-based indexing and operations.

```

# Convert the data frame to xts object
trend1_xts <- xts(trend1_data$AI, order.by = trend1_data$Week)

# View the first few rows of xts object
head(trend1_xts)

```

```

      [,1]
2019-05-26  14
2019-06-02  14
2019-06-09  14
2019-06-16  14
2019-06-23  14
2019-06-30  14

```

Visualization

The time series data is visualized to observe trends, seasonal patterns, and any apparent anomalies in the 'AI' search trends over the last five years.

```

plot(trend1_xts,
     type = "l",
     col = "darkred",
     lwd = 3,
     main = "'AI' Search Trends - Last 5 years")

```



DETRENDING AND STATIONARITY

In this section the stationarity assumption is addressed, since it is crucial in time series modeling. Non-stationary data can lead to unreliable and spurious results.

Transformations

Log transformation is applied to stabilize the variance.

```
# Log transformation
log_trend1_xts <- log(trend1_xts)
plot(log_trend1_xts,
     type = "l",
```

```
col = "darkred",
lwd = 3,
main = "'AI' Search Trends - Last 5 years - Log")
```

'AI' Search Trends – Last 5 years – Log

2019-05-26 / 2024-05-26



Differencing


Differencing is used to remove trends and achieve stationarity.

```
diff_log_trend1_xts <- diff(log_trend1_xts, differences = 1)
plot(diff_log_trend1_xts,
     type = "l",
     col = "darkred",
     lwd = 3,
```

```
main = "'AI' Search Trends - Last 5 years - Log")
```

'AI' Search Trends – Last 5 years – Log

2019-05-26 / 2024-05-26



The graph suggests that stationarity has been achieved but it has to be checked through the relative tests.

ADF Test for stationarity

The Augmented Dickey-Fuller (ADF) test checks for the presence of a unit root in the series, which indicates non-stationarity.

```
# Apply the ADF test to the differenced series
diff_log_trend1_xts <- diff_log_trend1_xts[-1,]
adf_test <- ur.df(diff_log_trend1_xts, type = "drift", lags = 0)
```

```
summary(adf_test)
```

```
#####
# Augmented Dickey-Fuller Test Unit Root Test #
#####

Test regression drift

Call:
lm(formula = z.diff ~ z.lag.1 + 1)

Residuals:
    Min       1Q   Median       3Q      Max
-0.16257 -0.03395 -0.00875  0.04063  0.36327

Coefficients:
                Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.008746    0.003855   2.269   0.0241 *
z.lag.1      -1.156534    0.061490 -18.809  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.06171 on 258 degrees of freedom
Multiple R-squared:  0.5783,    Adjusted R-squared:  0.5766
F-statistic: 353.8 on 1 and 258 DF,  p-value: < 2.2e-16

Value of test-statistic is: -18.8086 176.8808

Critical values for test statistics:
      1pct  5pct 10pct
tau2 -3.44 -2.87 -2.57
phi1  6.47  4.61  3.79
```

The test statistic value (-18.8086) is far below the critical values at all common significance levels (1%, 5%, and 10%), strongly indicating the rejection of the null hypothesis of a unit root. This result confirms that the time series is stationary.

PP Test for stationarity

The Phillips-Perron (PP) test also checks for a unit root, accounting for serial correlation in the error terms.

```
pp.test <- ur.pp(diff_log_trend1_xts, # tested series
                 type = c("Z-tau"),  # standardization of the test
                 ↪ statistic needed
                 model = c("constant")) # constant deterministic
                 ↪ component
# which means we assume that any trends in the data are stochastic
summary(pp.test)
```

```
#####
# Phillips-Perron Unit Root Test #
#####
```

Test regression with intercept

Call:

```
lm(formula = y ~ y.l1)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.16257	-0.03395	-0.00875	0.04063	0.36327

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.008746	0.003855	2.269	0.0241 *
y.l1	-0.156534	0.061490	-2.546	0.0115 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.06171 on 258 degrees of freedom

Multiple R-squared: 0.0245, Adjusted R-squared: 0.02072

F-statistic: 6.481 on 1 and 258 DF, p-value: 0.01149

Value of test-statistic, type: Z-tau is: -19.3793


```
      aux. Z statistics
Z-tau-mu      2.3374
```

Critical values for Z statistics:

```
      1pct      5pct      10pct
critical values -3.457006 -2.872754 -2.572697
```

The Z-tau test statistic value (-19.3793) is far below the critical values at all common significance levels (1%, 5%, and 10%), strongly indicating the rejection of the null hypothesis of a unit root. This confirms the results of the ADF test and we can say that the series is stationary.

KPSS Test

The KPSS is like the other 2 tests but the hypothesis are inverted.

```
kpss.test <- ur.kpss(diff_log_trend1_xts,
                     type = c("mu")) # constant deterministic
summary(kpss.test)
```

```
#####
# KPSS Unit Root Test #
#####
```

Test is of type: mu with 5 lags.

Value of test-statistic is: 0.5171

Critical value for a significance level of:

```
      10pct  5pct 2.5pct  1pct
critical values 0.347 0.463 0.574 0.739
```

In this case the test-statistic of **0.51** is lower than than 2.5% critical value (0.574) so we **cannot reject the null** about **stationarity** of the first differences at **2.5% and 1% significance level**. (however we reject it at 5% level).


By looking at the graph and from the results of the tests we can conclude that $AI \sim I(1)$.

ACF and PACF plots

ACF shows the correlation between the time series and its lagged versions, while PACF gives the partial correlation of a stationary time series with its own lagged values. Their plots help in identifying the appropriate AR and MA terms for ARIMA modeling.

```
# ACF and PACF plots of the differenced series
acf_plot <- ggAcf(diff_log_trend1_xts, main = "ACF of Differenced
  ↳ Log-Transformed Data")
pacf_plot <- ggPacf(diff_log_trend1_xts, main = "PACF of Differenced
  ↳ Log-Transformed Data")

gridExtra::grid.arrange(acf_plot, pacf_plot, ncol = 2)
```



As shown by the plots, we can see some significant spikes at lag 1 and 2 for the ACF plot and some at lag 1, 2 and 9 in the PACF plot.

MODELING

A grid search is conducted to identify the optimal ARIMA parameters based on AIC and BIC values. The results help in selecting the best combination of p, d, and q.

Grid search for ARIMA parameters

```
# Define the parameter grid
p_values <- c(1, 2, 9)
d_values <- c(1)
q_values <- c(1, 2, 8)

# Initialize an empty list to store results
results <- list()

# Perform grid search WITHOUT CONSTANT
for (p in p_values) {
  for (d in d_values) {
    for (q in q_values) {
      # Fit the ARIMA model
      model <- tryCatch(
        {
          Arima(log_trend1_xts, order = c(p, d, q))
        },
        error = function(e) NULL
      )

      # Check if model fitting was successful
      if (!is.null(model)) {
        # Extract AIC and BIC
        aic <- AIC(model)
        bic <- BIC(model)

        # Store the results
        results <- rbind(results, data.frame(p = p, d = d, q = q, AIC =
↪ aic, BIC = bic))
      }
    }
  }
}
```

```
# Convert results to a data frame
results_df <- as.data.frame(results)

# Print the results
print(results_df)
```

	p	d	q	AIC	BIC
1	1	1	1	-708.0181	-697.3245
2	1	1	2	-709.9838	-695.7257
3	1	1	8	-712.0973	-676.4521
4	2	1	1	-709.8640	-695.6060
5	2	1	2	-708.0825	-690.2599
6	2	1	8	-710.5610	-671.3512
7	9	1	1	-712.6472	-673.4375
8	9	1	2	-710.6835	-667.9092
9	9	1	8	-709.7928	-645.6314

Above we find the table with the respective AIC and BIC values for each model (without including a constant).

Now let's do the same thing by including a constant in the model and compare the results:

```
# Define the parameter grid
p_values <- c(1, 2, 9)
d_values <- c(1)
q_values <- c(1, 2, 8)

# Initialize an empty list to store results
results <- list()

# Perform grid search WITHOUT CONSTANT
for (p in p_values) {
  for (d in d_values) {
    for (q in q_values) {
      # Fit the ARIMA model
      model <- tryCatch(
        {
          Arima(log_trend1_xts, order = c(p, d, q), include.constant =
            TRUE)
        },
```

```

    error = function(e) NULL
  )

  # Check if model fitting was successful
  if (!is.null(model)) {
    # Extract AIC and BIC
    aic <- AIC(model)
    bic <- BIC(model)

    # Store the results
    results <- rbind(results, data.frame(p = p, d = d, q = q, AIC =
↪ aic, BIC = bic))
  }
}
}

# Convert results to a data frame
results_df <- as.data.frame(results)

# Print the results
print(results_df)

```

	p	d	q	AIC	BIC
1	1	1	1	-713.2871	-699.0290
2	1	1	2	-715.0502	-697.2276
3	1	1	8	-714.1144	-674.9047
4	2	1	1	-713.9599	-696.1373
5	2	1	2	-713.0656	-691.6785
6	2	1	8	-712.8204	-670.0462
7	9	1	1	-713.1831	-670.4089
8	9	1	2	-711.1839	-664.8451

According to the table we can see that values are generally lower in this case, implying that drift should be included in the model.

Let's now compare the best models in terms of AIC / BIC, that are ARIMA (1,1,1) since it has the best BIC, ARIMA (1,1,2) since it has the best AIC and ARIMA (2,1,1) since it has good values for both.

Model diagnostics

In this section the residuals are explored to make sure they don't present autocorrelation (i.e. they are white noise).


```
# Fit three different ARIMA models
model1 <- Arima(log_trend1_xts, order = c(1, 1, 1), include.constant =
  ↪ TRUE)
model2 <- Arima(log_trend1_xts, order = c(1, 1, 2), include.constant =
  ↪ TRUE)
model3 <- Arima(log_trend1_xts, order = c(2, 1, 1), include.constant =
  ↪ TRUE)

# Extract residuals
residuals_model1 <- residuals(model1)
residuals_model2 <- residuals(model2)
residuals_model3 <- residuals(model3)
```

ARIMA models have been ran and their respective residuals extracted.

Let's see their plots:

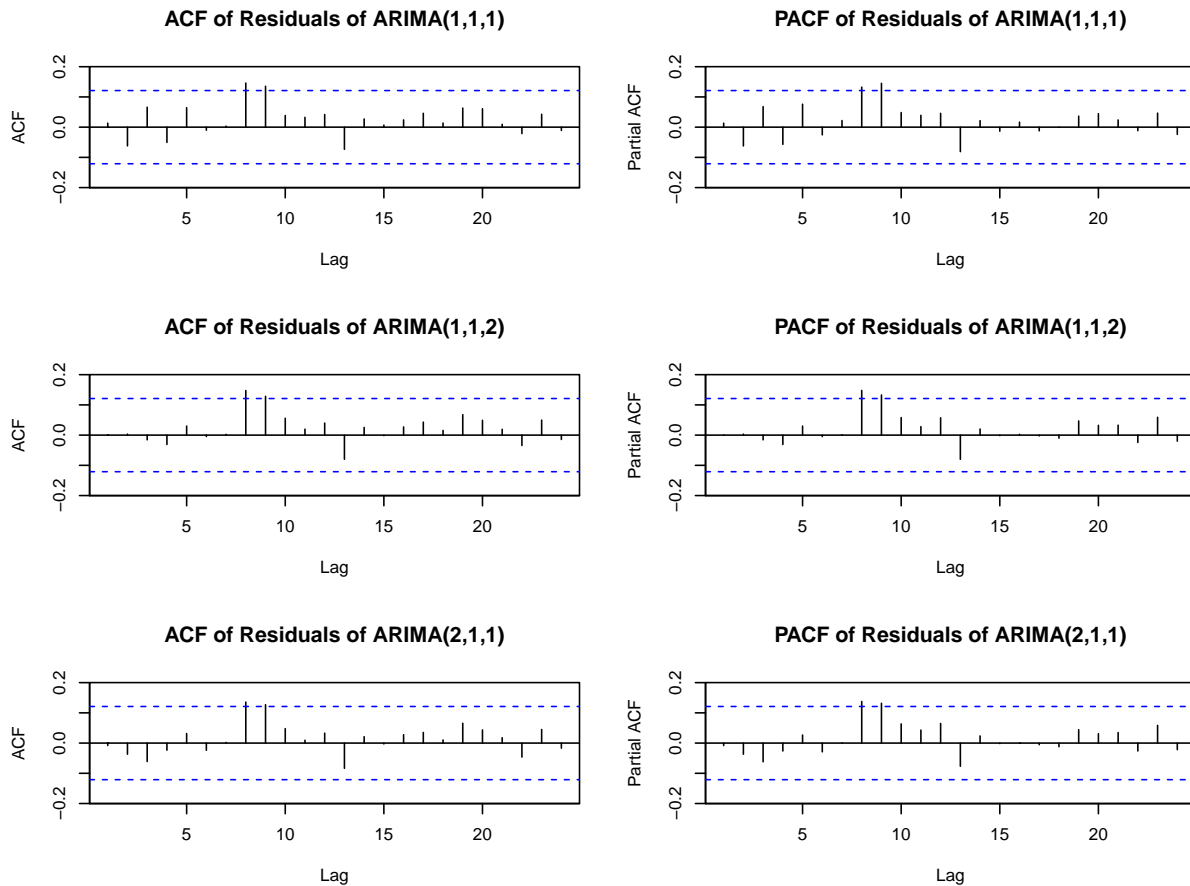
```
# Plot Residuals
par(mfrow = c(3, 2))
plot(residuals_model1, main = "Residuals of ARIMA(1,1,1)", ylab =
  ↪ "Residuals")
plot(residuals_model2, main = "Residuals of ARIMA(1,1,2)", ylab =
  ↪ "Residuals")
plot(residuals_model3, main = "Residuals of ARIMA(2,1,1)", ylab =
  ↪ "Residuals")
```



Residuals from different ARIMA models are plotted to visually inspect for white noise characteristics. Ideally, the residuals should appear random with no discernible patterns.

ACF and PACF plots of the residuals are also examined to check for any remaining autocorrelation:

```
# ACF and PACF of Residuals
par(mfrow = c(3, 2))
Acf(residuals_model1, main = "ACF of Residuals of ARIMA(1,1,1)")
Pacf(residuals_model1, main = "PACF of Residuals of ARIMA(1,1,1)")
Acf(residuals_model2, main = "ACF of Residuals of ARIMA(1,1,2)")
Pacf(residuals_model2, main = "PACF of Residuals of ARIMA(1,1,2)")
Acf(residuals_model3, main = "ACF of Residuals of ARIMA(2,1,1)")
Pacf(residuals_model3, main = "PACF of Residuals of ARIMA(2,1,1)")
```



For each model, we can see some spikes at lag 8 and 9, so to choose the best one between the three, some tests have to be performed:

- **Ljung-box test:** Checks the null hypothesis that the residuals are independently distributed (i.e. they are not autocorrelated).
- **Breusch-Godfrey test:** Checks the null hypothesis that there is no serial correlation in the residuals (i.e. they are white noise).

For each test, we would like the p-value to be greater than 0.05

We select number of **lags = 5** according to Ljung and Tsay ($\ln(T) = 5.5$)

```
# Perform tests on residuals to check for white noise
lb_test_model1 <- Box.test(residuals_model1, lag = 5, type =
  ↪ "Ljung-Box")
```



```

lb_test_model2 <- Box.test(residuals_model2, lag = 5, type =
  ↪ "Ljung-Box")
lb_test_model3 <- Box.test(residuals_model3, lag = 5, type =
  ↪ "Ljung-Box")
bg_test_model1 <- bgtest(residuals_model1 ~ fitted(model1), order = 5)
bg_test_model2 <- bgtest(residuals_model2 ~ fitted(model2), order = 5)
bg_test_model3 <- bgtest(residuals_model3 ~ fitted(model3), order = 5)

# Create a table to compare results
comparison_table <- data.frame(
  Model = c("ARIMA(1,1,1)", "ARIMA(1,1,2)", "ARIMA(2,1,1)"),
  AIC = c(AIC(model1), AIC(model2), AIC(model3)),
  BIC = c(BIC(model1), BIC(model2), BIC(model3)),
  LjungBox_p_value = c(lb_test_model1$p.value, lb_test_model2$p.value,
    ↪ lb_test_model3$p.value),
  BG_statistic = c(bg_test_model1$p.value, bg_test_model2$p.value,
    ↪ bg_test_model3$p.value)
)

# Print the comparison table
print(comparison_table)

```

	Model	AIC	BIC	LjungBox_p_value	BG_statistic
1	ARIMA(1,1,1)	-713.2871	-699.0290	0.5442502	0.4715572
2	ARIMA(1,1,2)	-715.0502	-697.2276	0.9894031	0.9781981
3	ARIMA(2,1,1)	-713.9599	-696.1373	0.8770163	0.7930028

What we can say from the comparison table is:

- For the Ljung-Box test, all three models (ARIMA(1,1,1), ARIMA(1,1,2), ARIMA(2,1,1)) have p-values well above 0.05 (0.5442502, 0.9894031, 0.8770163 respectively), suggesting that none of the models have significant autocorrelation in their residuals.
- For the BG test, p-values for the models are 0.4715572, 0.9781981, and 0.7930028 respectively. This suggests that for all of the three models the residuals are white noise.
- **ARIMA (1,1,2)** has the highest p-values for both tests and the best AIC value, so it is selected as the model that is going to be used for the forecasts.

Automatic model diagnostics

Now that the best model has been selected manually, some automatic model selection functions can be used and then compare the models' performances.

Auto AIC

Models with all combinations of p from 1 to 9 and q from 1 to 9 are analyzed and rated according to their AIC value:

```
arima.best.AIC <-  
  auto.arima(log_trend1_xts,  
             d = 1,           # parameter d of ARIMA model  
             max.p = 9,       # Maximum value of p  
             max.q = 9,       # Maximum value of q  
             max.order = 18,  # maximum p+q  
             start.p = 1,     # Starting value of p in stepwise  
             ↪ procedure  
             start.q = 1,     # Starting value of q in stepwise  
             ↪ procedure  
             ic = "aic",      # Information criterion to be used in  
             ↪ model selection.  
             stepwise = FALSE, # if FALSE considers all models  
             allowdrift = TRUE, # include a constant  
             trace = TRUE)    # show summary of all models considered
```

Fitting models using approximations to speed things up...

ARIMA(0,1,0)	: -698.4329
ARIMA(0,1,0) with drift	: -700.243
ARIMA(0,1,1)	: -703.1706
ARIMA(0,1,1) with drift	: -707.2513
ARIMA(0,1,2)	: -703.8547
ARIMA(0,1,2) with drift	: -709.4507
ARIMA(0,1,3)	: -703.6541
ARIMA(0,1,3) with drift	: -708.457
ARIMA(0,1,4)	: -702.1104
ARIMA(0,1,4) with drift	: -707.312
ARIMA(0,1,5)	: -701.705
ARIMA(0,1,5) with drift	: -706.5377

ARIMA(0,1,6)	: -699.9709
ARIMA(0,1,6) with drift	: -704.6477
ARIMA(0,1,7)	: -699.218
ARIMA(0,1,7) with drift	: -703.4153
ARIMA(0,1,8)	: -707.3909
ARIMA(0,1,8) with drift	: -709.9724
ARIMA(0,1,9)	: -707.22
ARIMA(0,1,9) with drift	: -709.1223
ARIMA(1,1,0)	: -700.5758
ARIMA(1,1,0) with drift	: -703.7308
ARIMA(1,1,1)	: -701.511
ARIMA(1,1,1) with drift	: -706.8663
ARIMA(1,1,2)	: -703.5201
ARIMA(1,1,2) with drift	: -708.5851
ARIMA(1,1,3)	: -701.5996
ARIMA(1,1,3) with drift	: -706.5911
ARIMA(1,1,4)	: -699.6244
ARIMA(1,1,4) with drift	: -704.8335
ARIMA(1,1,5)	: -708.2386
ARIMA(1,1,5) with drift	: -708.8586
ARIMA(1,1,6)	: -706.3513
ARIMA(1,1,6) with drift	: -706.9195
ARIMA(1,1,7)	: -704.7301
ARIMA(1,1,7) with drift	: -705.4502
ARIMA(1,1,8)	: -705.955
ARIMA(1,1,8) with drift	: -707.8649
ARIMA(1,1,9)	: -704.2214
ARIMA(1,1,9) with drift	: -706.116
ARIMA(2,1,0)	: -702.6465
ARIMA(2,1,0) with drift	: -707.5998
ARIMA(2,1,1)	: -702.3992
ARIMA(2,1,1) with drift	: -706.5471
ARIMA(2,1,2)	: -700.6169
ARIMA(2,1,2) with drift	: -705.696
ARIMA(2,1,3)	: -706.728
ARIMA(2,1,3) with drift	: -703.738
ARIMA(2,1,4)	: -705.2976
ARIMA(2,1,4) with drift	: -703.1931
ARIMA(2,1,5)	: -705.4142
ARIMA(2,1,5) with drift	: -706.0213
ARIMA(2,1,6)	: Inf
ARIMA(2,1,6) with drift	: Inf
ARIMA(2,1,7)	: Inf

ARIMA(2,1,7) with drift	: -702.5663
ARIMA(2,1,8)	: -703.4342
ARIMA(2,1,8) with drift	: -705.6162
ARIMA(2,1,9)	: Inf
ARIMA(2,1,9) with drift	: -704.1108
ARIMA(3,1,0)	: -700.0285
ARIMA(3,1,0) with drift	: -704.6663
ARIMA(3,1,1)	: -699.5123
ARIMA(3,1,1) with drift	: -703.8845
ARIMA(3,1,2)	: -697.6481
ARIMA(3,1,2) with drift	: -702.8165
ARIMA(3,1,3)	: -705.4334
ARIMA(3,1,3) with drift	: Inf
ARIMA(3,1,4)	: Inf
ARIMA(3,1,4) with drift	: -702.0234
ARIMA(3,1,5)	: Inf
ARIMA(3,1,5) with drift	: Inf
ARIMA(3,1,6)	: Inf
ARIMA(3,1,6) with drift	: Inf
ARIMA(3,1,7)	: Inf
ARIMA(3,1,7) with drift	: Inf
ARIMA(3,1,8)	: -700.5921
ARIMA(3,1,8) with drift	: -702.8043
ARIMA(3,1,9)	: -698.7332
ARIMA(3,1,9) with drift	: -701.695
ARIMA(4,1,0)	: -698.0541
ARIMA(4,1,0) with drift	: -703.78
ARIMA(4,1,1)	: -697.0899
ARIMA(4,1,1) with drift	: -702.6024
ARIMA(4,1,2)	: Inf
ARIMA(4,1,2) with drift	: Inf
ARIMA(4,1,3)	: Inf
ARIMA(4,1,3) with drift	: Inf
ARIMA(4,1,4)	: -707.0912
ARIMA(4,1,4) with drift	: -707.793
ARIMA(4,1,5)	: Inf
ARIMA(4,1,5) with drift	: Inf
ARIMA(4,1,6)	: Inf
ARIMA(4,1,6) with drift	: Inf
ARIMA(4,1,7)	: Inf
ARIMA(4,1,7) with drift	: Inf
ARIMA(4,1,8)	: Inf
ARIMA(4,1,8) with drift	: Inf

ARIMA(4,1,9)	: Inf
ARIMA(4,1,9) with drift	: Inf
ARIMA(5,1,0)	: -696.6702
ARIMA(5,1,0) with drift	: -701.506
ARIMA(5,1,1)	: -694.7047
ARIMA(5,1,1) with drift	: -699.5596
ARIMA(5,1,2)	: -700.9894
ARIMA(5,1,2) with drift	: Inf
ARIMA(5,1,3)	: Inf
ARIMA(5,1,3) with drift	: Inf
ARIMA(5,1,4)	: Inf
ARIMA(5,1,4) with drift	: Inf
ARIMA(5,1,5)	: Inf
ARIMA(5,1,5) with drift	: Inf
ARIMA(5,1,6)	: Inf
ARIMA(5,1,6) with drift	: Inf
ARIMA(5,1,7)	: Inf
ARIMA(5,1,7) with drift	: Inf
ARIMA(5,1,8)	: -695.2927
ARIMA(5,1,8) with drift	: -698.1751
ARIMA(5,1,9)	: -693.2953
ARIMA(5,1,9) with drift	: -696.3847
ARIMA(6,1,0)	: -695.1703
ARIMA(6,1,0) with drift	: -700.8807
ARIMA(6,1,1)	: -693.1896
ARIMA(6,1,1) with drift	: -699.0504
ARIMA(6,1,2)	: -696.2897
ARIMA(6,1,2) with drift	: -697.785
ARIMA(6,1,3)	: Inf
ARIMA(6,1,3) with drift	: Inf
ARIMA(6,1,4)	: Inf
ARIMA(6,1,4) with drift	: -698.1418
ARIMA(6,1,5)	: Inf
ARIMA(6,1,5) with drift	: Inf
ARIMA(6,1,6)	: Inf
ARIMA(6,1,6) with drift	: Inf
ARIMA(6,1,7)	: Inf
ARIMA(6,1,7) with drift	: Inf
ARIMA(6,1,8)	: -694.501
ARIMA(6,1,8) with drift	: -697.8074
ARIMA(6,1,9)	: Inf
ARIMA(6,1,9) with drift	: Inf
ARIMA(7,1,0)	: -692.1834

ARIMA(7,1,0) with drift	: -698.1525
ARIMA(7,1,1)	: -700.4668
ARIMA(7,1,1) with drift	: -696.2813
ARIMA(7,1,2)	: -701.3723
ARIMA(7,1,2) with drift	: -698.7333
ARIMA(7,1,3)	: -699.6977
ARIMA(7,1,3) with drift	: Inf
ARIMA(7,1,4)	: Inf
ARIMA(7,1,4) with drift	: Inf
ARIMA(7,1,5)	: Inf
ARIMA(7,1,5) with drift	: Inf
ARIMA(7,1,6)	: Inf
ARIMA(7,1,6) with drift	: Inf
ARIMA(7,1,7)	: Inf
ARIMA(7,1,7) with drift	: Inf
ARIMA(7,1,8)	: Inf
ARIMA(7,1,8) with drift	: Inf
ARIMA(7,1,9)	: Inf
ARIMA(7,1,9) with drift	: Inf
ARIMA(8,1,0)	: -692.9176
ARIMA(8,1,0) with drift	: -697.6092
ARIMA(8,1,1)	: -701.8554
ARIMA(8,1,1) with drift	: -702.9817
ARIMA(8,1,2)	: -700.8535
ARIMA(8,1,2) with drift	: -702.6568
ARIMA(8,1,3)	: -698.8542
ARIMA(8,1,3) with drift	: -700.6602
ARIMA(8,1,4)	: -697.0867
ARIMA(8,1,4) with drift	: -699.2619
ARIMA(8,1,5)	: -695.2056
ARIMA(8,1,5) with drift	: Inf
ARIMA(8,1,6)	: -694.4045
ARIMA(8,1,6) with drift	: -695.6406
ARIMA(8,1,7)	: Inf
ARIMA(8,1,7) with drift	: Inf
ARIMA(8,1,8)	: Inf
ARIMA(8,1,8) with drift	: Inf
ARIMA(8,1,9)	: Inf
ARIMA(8,1,9) with drift	: Inf
ARIMA(9,1,0)	: -697.4957
ARIMA(9,1,0) with drift	: -700.5683
ARIMA(9,1,1)	: -700.1773
ARIMA(9,1,1) with drift	: -701.5039

```

ARIMA(9,1,2) : -698.199
ARIMA(9,1,2) with drift : Inf
ARIMA(9,1,3) : -696.1991
ARIMA(9,1,3) with drift : -697.6981
ARIMA(9,1,4) : -694.9972
ARIMA(9,1,4) with drift : -696.5029
ARIMA(9,1,5) : -694.1092
ARIMA(9,1,5) with drift : Inf
ARIMA(9,1,6) : Inf
ARIMA(9,1,6) with drift : Inf
ARIMA(9,1,7) : Inf
ARIMA(9,1,7) with drift : Inf
ARIMA(9,1,8) : Inf
ARIMA(9,1,8) with drift : Inf
ARIMA(9,1,9) : Inf
ARIMA(9,1,9) with drift : Inf

```

Now re-fitting the best model(s) without approximations...

Best model: ARIMA(0,1,8) with drift

According to the automatic function, best model is ARIMA(0,1,8) with drift.

Let's test the significance of its coefficients:

```
coeftest(arima.best.AIC)
```

z test of coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
ma1	-0.2106522	0.0611543	-3.4446	0.0005719	***
ma2	-0.1520129	0.0623873	-2.4366	0.0148261	*
ma3	0.0696407	0.0648367	1.0741	0.2827809	
ma4	-0.0788677	0.0634919	-1.2422	0.2141736	
ma5	0.1021360	0.0641732	1.5916	0.1114821	
ma6	-0.0258656	0.0641125	-0.4034	0.6866237	
ma7	0.0051393	0.0613435	0.0838	0.9332318	
ma8	0.1854018	0.0642254	2.8867	0.0038926	**

```
drift 0.0074155 0.0032582 2.2759 0.0228510 *
```

```
---
```


```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- ma1, ma2, and ma8 are significant while ma3 to ma7 are not and might indicate some over-parametrization. Moreover, the drift term is also significant, suggesting it should be included in the model.


Model diagnostics

Let's again look at the ACF and PACF plots:

```
acf(resid(arima.best.AIC),  
    lag.max = 48,  
    lwd = 7,  
    col = "darkred",  
    na.action = na.pass,  
    ylim = c(-0.08, 0.05))
```

```
pacf(resid(arima.best.AIC),  
      lag.max = 48,  
      lwd = 7,  
      col = "darkred",  
      na.action = na.pass)
```



Residuals look overall good although showing 2 significant spikes, so let's check the statistics.

Performing the Ljung-Box test and the BG test for the residuals:

```
Box.test(resid(arima.best.AIC), type = "Ljung-Box", lag = 5)
```

Box-Ljung test

```
data: resid(arima.best.AIC)
X-squared = 0.14463, df = 5, p-value = 0.9996
```

```
bgtest(arima.best.AIC$residuals ~ fitted(arima.best.AIC), order = 5)
```

Breusch-Godfrey test for serial correlation of order up to 5

```
data: arima.best.AIC$residuals ~ fitted(arima.best.AIC)
LM test = 0.20204, df = 5, p-value = 0.9991
```

For both tests the p.value is 0.999 meaning that residuals are not autocorrelated and they are white noise.

Let's now compare AIC and BIC values for both the automatic and the manual model:

```
AIC(arima.best.AIC, model2)
```

	df	AIC
arima.best.AIC	10	-715.1920
model2	5	-715.0502

```
BIC(arima.best.AIC, model2)
```

	df	BIC
arima.best.AIC	10	-679.5468
model2	5	-697.2276

As the table shows, the automatic model is just slightly better in terms of AIC but the manual model is way better in BIC.

Auto BIC

Let's now do the same procedure by selecting the best model in terms of BIC:

```
arima.best.BIC <-
  auto.arima(log_trend1_xts,
    d = 1,           # parameter d of ARIMA model
    max.p = 9,       # Maximum value of p
    max.q = 9,       # Maximum value of q
    max.order = 18,  # maximum p+q
    start.p = 1,     # Starting value of p in stepwise
    ↪ procedure
    start.q = 1,     # Starting value of q in stepwise
    ↪ procedure
```

```

ic = "bic",          # Information criterion to be used in
  ↪ model selection.
stepwise = FALSE,    # if FALSE considers all models
allowdrift = TRUE,    # include a constant
trace = TRUE)        # show summary of all models considered

```

Fitting models using approximations to speed things up...

```

ARIMA(0,1,0)           : -694.8684
ARIMA(0,1,0) with drift : -693.114
ARIMA(0,1,1)           : -696.0416
ARIMA(0,1,1) with drift : -696.5577
ARIMA(0,1,2)           : -693.1612
ARIMA(0,1,2) with drift : -695.1927
ARIMA(0,1,3)           : -689.396
ARIMA(0,1,3) with drift : -690.6344
ARIMA(0,1,4)           : -684.2878
ARIMA(0,1,4) with drift : -685.9249
ARIMA(0,1,5)           : -680.3179
ARIMA(0,1,5) with drift : -681.586
ARIMA(0,1,6)           : -675.0193
ARIMA(0,1,6) with drift : -676.1316
ARIMA(0,1,7)           : -670.7019
ARIMA(0,1,7) with drift : -671.3346
ARIMA(0,1,8)           : -675.3102
ARIMA(0,1,8) with drift : -674.3272
ARIMA(0,1,9)           : -671.5748
ARIMA(0,1,9) with drift : -669.9126
ARIMA(1,1,0)           : -693.4467
ARIMA(1,1,0) with drift : -693.0372
ARIMA(1,1,1)           : -690.8175
ARIMA(1,1,1) with drift : -692.6082
ARIMA(1,1,2)           : -689.262
ARIMA(1,1,2) with drift : -690.7625
ARIMA(1,1,3)           : -683.777
ARIMA(1,1,3) with drift : -685.204
ARIMA(1,1,4)           : -678.2373
ARIMA(1,1,4) with drift : -679.8819
ARIMA(1,1,5)           : -683.2869
ARIMA(1,1,5) with drift : -680.3425

```

ARIMA(1,1,6)	: -677.8351
ARIMA(1,1,6) with drift	: -674.8388
ARIMA(1,1,7)	: -672.6494
ARIMA(1,1,7) with drift	: -669.805
ARIMA(1,1,8)	: -670.3098
ARIMA(1,1,8) with drift	: -668.6551
ARIMA(1,1,9)	: -665.0116
ARIMA(1,1,9) with drift	: -663.3417
ARIMA(2,1,0)	: -691.9529
ARIMA(2,1,0) with drift	: -693.3418
ARIMA(2,1,1)	: -688.1411
ARIMA(2,1,1) with drift	: -688.7245
ARIMA(2,1,2)	: -682.7943
ARIMA(2,1,2) with drift	: -684.3088
ARIMA(2,1,3)	: -685.3409
ARIMA(2,1,3) with drift	: -678.7864
ARIMA(2,1,4)	: -680.3459
ARIMA(2,1,4) with drift	: -674.6769
ARIMA(2,1,5)	: -676.8981
ARIMA(2,1,5) with drift	: -673.9406
ARIMA(2,1,6)	: Inf
ARIMA(2,1,6) with drift	: Inf
ARIMA(2,1,7)	: Inf
ARIMA(2,1,7) with drift	: -663.3566
ARIMA(2,1,8)	: -664.2245
ARIMA(2,1,8) with drift	: -662.842
ARIMA(2,1,9)	: Inf
ARIMA(2,1,9) with drift	: -657.772
ARIMA(3,1,0)	: -685.7705
ARIMA(3,1,0) with drift	: -686.8437
ARIMA(3,1,1)	: -681.6897
ARIMA(3,1,1) with drift	: -682.4974
ARIMA(3,1,2)	: -676.261
ARIMA(3,1,2) with drift	: -677.8648
ARIMA(3,1,3)	: -680.4817
ARIMA(3,1,3) with drift	: Inf
ARIMA(3,1,4)	: Inf
ARIMA(3,1,4) with drift	: -669.9427
ARIMA(3,1,5)	: Inf
ARIMA(3,1,5) with drift	: Inf
ARIMA(3,1,6)	: Inf
ARIMA(3,1,6) with drift	: Inf
ARIMA(3,1,7)	: Inf

ARIMA(3,1,7) with drift	: Inf
ARIMA(3,1,8)	: -657.8178
ARIMA(3,1,8) with drift	: -656.4656
ARIMA(3,1,9)	: -652.3944
ARIMA(3,1,9) with drift	: -651.7917
ARIMA(4,1,0)	: -680.2315
ARIMA(4,1,0) with drift	: -682.3929
ARIMA(4,1,1)	: -675.7027
ARIMA(4,1,1) with drift	: -677.6508
ARIMA(4,1,2)	: Inf
ARIMA(4,1,2) with drift	: Inf
ARIMA(4,1,3)	: Inf
ARIMA(4,1,3) with drift	: Inf
ARIMA(4,1,4)	: -675.0105
ARIMA(4,1,4) with drift	: -672.1478
ARIMA(4,1,5)	: Inf
ARIMA(4,1,5) with drift	: Inf
ARIMA(4,1,6)	: Inf
ARIMA(4,1,6) with drift	: Inf
ARIMA(4,1,7)	: Inf
ARIMA(4,1,7) with drift	: Inf
ARIMA(4,1,8)	: Inf
ARIMA(4,1,8) with drift	: Inf
ARIMA(4,1,9)	: Inf
ARIMA(4,1,9) with drift	: Inf
ARIMA(5,1,0)	: -675.2831
ARIMA(5,1,0) with drift	: -676.5543
ARIMA(5,1,1)	: -669.7531
ARIMA(5,1,1) with drift	: -671.0435
ARIMA(5,1,2)	: -672.4733
ARIMA(5,1,2) with drift	: Inf
ARIMA(5,1,3)	: Inf
ARIMA(5,1,3) with drift	: Inf
ARIMA(5,1,4)	: Inf
ARIMA(5,1,4) with drift	: Inf
ARIMA(5,1,5)	: Inf
ARIMA(5,1,5) with drift	: Inf
ARIMA(5,1,6)	: Inf
ARIMA(5,1,6) with drift	: Inf
ARIMA(5,1,7)	: Inf
ARIMA(5,1,7) with drift	: Inf
ARIMA(5,1,8)	: -645.3894
ARIMA(5,1,8) with drift	: -644.7073

ARIMA(5,1,9)	: -639.8275
ARIMA(5,1,9) with drift	: -639.3524
ARIMA(6,1,0)	: -670.2186
ARIMA(6,1,0) with drift	: -672.3645
ARIMA(6,1,1)	: -664.6734
ARIMA(6,1,1) with drift	: -666.9697
ARIMA(6,1,2)	: -664.209
ARIMA(6,1,2) with drift	: -662.1398
ARIMA(6,1,3)	: Inf
ARIMA(6,1,3) with drift	: Inf
ARIMA(6,1,4)	: Inf
ARIMA(6,1,4) with drift	: -655.3676
ARIMA(6,1,5)	: Inf
ARIMA(6,1,5) with drift	: Inf
ARIMA(6,1,6)	: Inf
ARIMA(6,1,6) with drift	: Inf
ARIMA(6,1,7)	: Inf
ARIMA(6,1,7) with drift	: Inf
ARIMA(6,1,8)	: -641.0332
ARIMA(6,1,8) with drift	: -640.7751
ARIMA(6,1,9)	: Inf
ARIMA(6,1,9) with drift	: Inf
ARIMA(7,1,0)	: -663.6672
ARIMA(7,1,0) with drift	: -666.0718
ARIMA(7,1,1)	: -668.3861
ARIMA(7,1,1) with drift	: -660.6361
ARIMA(7,1,2)	: -665.7271
ARIMA(7,1,2) with drift	: -659.5236
ARIMA(7,1,3)	: -660.488
ARIMA(7,1,3) with drift	: Inf
ARIMA(7,1,4)	: Inf
ARIMA(7,1,4) with drift	: Inf
ARIMA(7,1,5)	: Inf
ARIMA(7,1,5) with drift	: Inf
ARIMA(7,1,6)	: Inf
ARIMA(7,1,6) with drift	: Inf
ARIMA(7,1,7)	: Inf
ARIMA(7,1,7) with drift	: Inf
ARIMA(7,1,8)	: Inf
ARIMA(7,1,8) with drift	: Inf
ARIMA(7,1,9)	: Inf
ARIMA(7,1,9) with drift	: Inf
ARIMA(8,1,0)	: -660.8369

ARIMA(8,1,0) with drift	: -661.964
ARIMA(8,1,1)	: -666.2102
ARIMA(8,1,1) with drift	: -663.772
ARIMA(8,1,2)	: -661.6438
ARIMA(8,1,2) with drift	: -659.8825
ARIMA(8,1,3)	: -656.0799
ARIMA(8,1,3) with drift	: -654.3215
ARIMA(8,1,4)	: -650.7479
ARIMA(8,1,4) with drift	: -649.3587
ARIMA(8,1,5)	: -645.3023
ARIMA(8,1,5) with drift	: Inf
ARIMA(8,1,6)	: -640.9367
ARIMA(8,1,6) with drift	: -638.6083
ARIMA(8,1,7)	: Inf
ARIMA(8,1,7) with drift	: Inf
ARIMA(8,1,8)	: Inf
ARIMA(8,1,8) with drift	: Inf
ARIMA(8,1,9)	: Inf
ARIMA(8,1,9) with drift	: Inf
ARIMA(9,1,0)	: -661.8505
ARIMA(9,1,0) with drift	: -661.3586
ARIMA(9,1,1)	: -660.9676
ARIMA(9,1,1) with drift	: -658.7296
ARIMA(9,1,2)	: -655.4248
ARIMA(9,1,2) with drift	: Inf
ARIMA(9,1,3)	: -649.8604
ARIMA(9,1,3) with drift	: -647.7948
ARIMA(9,1,4)	: -645.0939
ARIMA(9,1,4) with drift	: -643.0351
ARIMA(9,1,5)	: -640.6414
ARIMA(9,1,5) with drift	: Inf
ARIMA(9,1,6)	: Inf
ARIMA(9,1,6) with drift	: Inf
ARIMA(9,1,7)	: Inf
ARIMA(9,1,7) with drift	: Inf
ARIMA(9,1,8)	: Inf
ARIMA(9,1,8) with drift	: Inf
ARIMA(9,1,9)	: Inf
ARIMA(9,1,9) with drift	: Inf

Now re-fitting the best model(s) without approximations...

Best model: ARIMA(0,1,1) with drift

In this case ARIMA(0,1,1) with drift is selected.

Diagnostics

Significancy of the coefficients is tested:

```
coeftest(arima.best.BIC)
```


z test of coefficients:

```
      Estimate Std. Error z value Pr(>|z|)
ma1   -0.2134061  0.0685149 -3.1147 0.001841 **
drift   0.0075367  0.0029766  2.5320 0.011341 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

As shown by the results, both the ma1 and the drift are significant at 5% level.

Let's now inspect the ACF and PACF plots:

```
par(mfrow=c(1,1))
acf(resid(arima.best.BIC),
    lag.max = 48,
    lwd = 7,
    col = "darkred",
    na.action = na.pass,
    ylim = c(-0.06, 0.06))
```

For this model way more spikes are observable, suggesting a check on the autocorrelation of the residuals:

```
Box.test(resid(arima.best.BIC), type = "Ljung-Box", lag = 5)
```

Box-Ljung test

```
data: resid(arima.best.BIC)
X-squared = 5.872, df = 5, p-value = 0.3189
```

```
bgtest(arima.best.BIC$residuals ~ fitted(arima.best.BIC), order = 5)
```

Breusch-Godfrey test for serial correlation of order up to 5

```
data: arima.best.BIC$residuals ~ fitted(arima.best.BIC)
LM test = 7.4087, df = 5, p-value = 0.192
```

According to the tests there is no significant autocorrelation in the residuals, even though the values are pretty low compared to the other models, suggesting there might be some issues.

Let's also check its values of BIC/AIC compared with the other models:

```
BIC(arima.best.BIC, arima.best.AIC, model2)
```

	df	BIC
arima.best.BIC	3	-702.0548
arima.best.AIC	10	-679.5468
model2	5	-697.2276

```
AIC(arima.best.BIC, arima.best.AIC, model2)
```

	df	AIC
arima.best.BIC	3	-712.7484
arima.best.AIC	10	-715.1920
model2	5	-715.0502

Also in this case, the values of BIC is higher for the automatic model but the value of AIC is higher for the manual one.

FORECASTING

This section is going to be about forecasting, where the performance of predicting the next values of two different models are compared.

Manual model

First of all, data is split in an “in sample” period and an “out-of-sample” period that is going to be used to test the model performances.

```
# Define the split point
split_point <- floor(0.9 * length(log_trend1_xts))

# Split the data into training and testing sets
train_data <- window(log_trend1_xts, end =
  ↪ index(log_trend1_xts)[split_point])
test_data <- window(log_trend1_xts, start =
  ↪ index(log_trend1_xts)[split_point + 1])
```

Then, the model is ran on the in-sample period:

```
manual_model <- Arima(train_data, order = c(1, 1, 2), include.constant =
  ↪ TRUE)
```

Let's now forecast the values for the out-of-sample period:


```
# Generate out-of-sample forecasts for the length of the test set
forecast_horizon <- length(test_data)
out_of_sample_forecast <- forecast(manual_model, h = forecast_horizon)
forecasts_data <- data.frame(f_mean =
  ↪ as.numeric(out_of_sample_forecast$mean),
  f_lower =
    ↪ as.numeric(out_of_sample_forecast$lower[,
    ↪ 2]),
  f_upper =
    ↪ as.numeric(out_of_sample_forecast$upper[,
    ↪ 2]))
```

Now that the forecasts are generated, the data is manipulated a bit to make sure a graph can be shown:

```
AI.oos <- test_data
# Ensure names match for consistency
names(AI.oos) <- "V1"
colnames(train_data) <- "V1"
AI2 <- rbind(train_data[, "V1"], AI.oos)
forecasts_xts <- xts(forecasts_data, order.by = index(AI.oos))
forecasted_data <- merge(AI2, forecasts_xts)
```

Let's now see the graph of the time series with the forecasted values:

```
plot(forecasted_data[, c("V1", "f_mean", "f_lower", "f_upper")],
     major.ticks = "years",
     grid.ticks.on = "years",
     grid.ticks.lty = 3,
     main = "forecast of 'AI' search term",
     col = c("black", "blue", "red", "red"))
```



Just from the graph the predictions look pretty good, but some measures have to be analyzed.

The data has been transformed back to normal and error measures analyzed:

- **MAE** (Mean Absolute Error)
- **MSE** (Mean Square Error)
- **MAPE** (Mean Absolute Percentage Error)

- **MAPE** (Adjusted Mean Absolute Percentage Error)

```
eval_data_oos <- tail(forecasted_data, length(test_data))
eval_data_oos$V1 <- exp(eval_data_oos$V1)
eval_data_oos$f_mean <- exp(eval_data_oos$f_mean)
eval_data_oos$f_lower <- exp(eval_data_oos$f_lower)
eval_data_oos$f_upper <- exp(eval_data_oos$f_upper)
eval_data_oos$mae <- abs(eval_data_oos$V1 - eval_data_oos$f_mean)
eval_data_oos$mse <- (eval_data_oos$V1 - eval_data_oos$f_mean) ^ 2
eval_data_oos$mape <- abs((eval_data_oos$V1 -
  ↪ eval_data_oos$f_mean)/eval_data_oos$V1)
eval_data_oos$amape <- abs((eval_data_oos$V1 -
  ↪ eval_data_oos$f_mean)/(eval_data_oos$V1 + eval_data_oos$f_mean))
eval_data_oos
```

	V1	f_mean	f_lower	f_upper	mae	mse	mape
2023-11-26	83	82.68320	73.19259	93.40442	0.3167997	0.10036202	0.003816863
2023-12-03	85	82.68854	70.60995	96.83330	2.3114628	5.34286019	0.027193680
2023-12-10	87	83.66206	69.94611	100.06762	3.3379394	11.14183972	0.038367120
2023-12-17	75	84.10251	68.72737	102.91724	9.1025071	82.85563541	0.121366761
2023-12-24	77	84.84834	68.06671	105.76744	7.8483445	61.59651181	0.101926552
2023-12-31	80	85.43113	67.30648	108.43649	5.4311327	29.49720221	0.067889159
2024-01-07	88	86.11260	66.74871	111.09397	1.8874021	3.56228679	0.021447751
2024-01-14	89	86.74658	66.20154	113.66757	2.2534247	5.07792301	0.025319379
2024-01-21	85	87.41478	65.74648	116.22437	2.4147757	5.83114148	0.028409125
2024-01-28	85	88.07161	65.32361	118.74127	3.0716086	9.43477945	0.036136572
2024-02-04	82	88.74260	64.95325	121.24488	6.7426019	45.46268004	0.082226852
2024-02-11	91	89.41355	64.61488	123.72976	1.5864466	2.51681271	0.017433479
2024-02-18	97	90.09246	64.31209	126.20724	6.9075431	47.71415141	0.071211784
2024-02-25	93	90.77491	64.03675	128.67742	2.2250932	4.95103989	0.023925734
2024-03-03	91	91.46342	63.78821	131.14584	0.4634247	0.21476248	0.005092579
2024-03-10	92	92.15666	63.56257	133.61401	0.1566630	0.02454331	0.001702859
2024-03-17	95	92.85544	63.35838	136.08511	2.1445639	4.59915426	0.022574357
2024-03-24	90	93.55935	63.17336	138.56082	3.5593509	12.66897910	0.039548344
2024-03-31	88	94.26869	63.00611	141.04324	6.2686895	39.29646793	0.071235108
2024-04-07	90	94.98336	62.85510	143.53390	4.9833571	24.83384838	0.055370635
2024-04-14	92	95.70347	62.71918	146.03434	3.7034701	13.71569097	0.040255110
2024-04-21	92	96.42903	62.59721	148.54587	4.4290274	19.61628338	0.048141602
2024-04-28	92	97.16009	62.48826	151.06972	5.1600938	26.62656798	0.056087976
2024-05-05	91	97.89670	62.39145	153.60699	6.8966980	47.56444288	0.075787890
2024-05-12	94	98.63889	62.30603	156.15872	4.6388893	21.51929351	0.049349886
2024-05-19	100	99.38671	62.23131	158.72584	0.6132941	0.37612968	0.006132941

2024-05-26	100	100.14019	62.16667	161.30924	0.1401928	0.01965402	0.001401928
------------	-----	-----------	----------	-----------	-----------	------------	-------------

	amape
2023-11-26	0.0019120807
2023-12-03	0.0137842623
2023-12-10	0.0195587668
2023-12-17	0.0572115881
2023-12-24	0.0484919667
2023-12-31	0.0328301729
2024-01-07	0.0108401239
2024-01-14	0.0128220122
2024-01-21	0.0140056190
2024-01-28	0.0177476169
2024-02-04	0.0394898625
2024-02-11	0.0087933891
2024-02-18	0.0369204788
2024-02-25	0.0121077098
2024-03-03	0.0025398226
2024-03-10	0.0008507052
2024-03-17	0.0114160331
2024-03-24	0.0193907361
2024-03-31	0.0343925745
2024-04-07	0.0269394891
2024-04-14	0.0197304297
2024-04-21	0.0235050163
2024-04-28	0.0272789767
2024-05-05	0.0365104210
2024-05-12	0.0240807517
2024-05-19	0.0030759028
2024-05-26	0.0007004730

Let's save the average errors for the out-of-sample evaluation data:

```
perf_manual_model = colMeans(eval_data_oos[, c("mae", "mse", "mape",
↪ "amape")])
```

Values will be analyzed and compared to the ones of the automatic model.

Automatic model

As an automatic model to compare, ARIMA(0,1,8) is chosen due to its better behavior of the residuals and similar AIC/BIC values.


```
automatic_model <- Arima(train_data, order = c(0, 1, 8),
  ↪ include.constant = TRUE)
```

Let's now forecast the values for the out-of-sample period:

```
# Generate out-of-sample forecasts for the length of the test set
out_of_sample_forecast_auto <- forecast(automatic_model, h =
  ↪ forecast_horizon)

forecasts_data_auto <- data.frame(f_mean =
  ↪ as.numeric(out_of_sample_forecast_auto$mean),
                                f_lower =
  ↪ as.numeric(out_of_sample_forecast_auto$lower[,
  ↪ 2]),
                                f_upper =
  ↪ as.numeric(out_of_sample_forecast_auto$upper[,
  ↪ 2]))
```

Now that the forecasts are generated, the data is manipulated a bit to make sure a graph can be shown:

```
forecasts_auto_xts <- xts(forecasts_data_auto,
  order.by = index(AI.oos))
forecasted_data_auto <- merge(AI2, forecasts_auto_xts)
```

Let's now see the graph of the time series with the forecasted values by the automatic model:

```
plot(forecasted_data_auto[, c("V1", "f_mean", "f_lower", "f_upper")],
  major.ticks = "years",
  grid.ticks.on = "years",
  grid.ticks.lty = 3,
  main = "forecast of 'AI' search term",
  col = c("black", "blue", "red", "red"))
```



Also in this case, the prediction don't look far from the actual values, but we can't tell just by looking at the graph which of the two models is performing better.

So, again the error measures are computed and saved to be compared with the other model:

```
eval_data_oos_auto <- tail(forecasted_data_auto, length(test_data))
eval_data_oos_auto$V1 <- exp(eval_data_oos_auto$V1)
eval_data_oos_auto$f_mean <- exp(eval_data_oos_auto$f_mean)
eval_data_oos_auto$f_lower <- exp(eval_data_oos_auto$f_lower)
eval_data_oos_auto$f_upper <- exp(eval_data_oos_auto$f_upper)
eval_data_oos_auto$mae <- abs(eval_data_oos_auto$V1 -
  ↪ eval_data_oos_auto$f_mean)
eval_data_oos_auto$mse <- (eval_data_oos_auto$V1 -
  ↪ eval_data_oos_auto$f_mean) ^ 2
```

```
eval_data_oos_auto$mape <- abs((eval_data_oos_auto$V1 -
  ↪ eval_data_oos_auto$f_mean)/eval_data_oos_auto$V1)
eval_data_oos_auto$amape <- abs((eval_data_oos_auto$V1 -
  ↪ eval_data_oos_auto$f_mean)/(eval_data_oos_auto$V1 +
  ↪ eval_data_oos_auto$f_mean))

perf_auto_model = colMeans(eval_data_oos_auto[, c("mae", "mse", "mape",
  ↪ "amape")])
```

Performance comparison

Let's now compare the error measures of the two models to select the best one:

```
# Combine the performance metrics into a single data frame
combined_perf <- data.frame(
  ARIMA_1_1_2 = perf_manual_model,
  ARIMA_0_1_8 = perf_auto_model
)

# Print the combined table
print(combined_perf)
```

	ARIMA_1_1_2	ARIMA_0_1_8
mae	3.65165914	4.53356390
mse	19.48744607	31.70184775
mape	0.04219822	0.05237733
amape	0.02062693	0.02519301

- **Overall Performance:** The ARIMA(1,1,2) model outperforms the ARIMA(0,1,8) model in all error metrics (MAE, MSE, MAPE, and AMAPE).
- **Model Choice:** Based on these error measures, the ARIMA(1,1,2) model appears to be a better choice for forecasting, as it consistently provides more accurate predictions with smaller errors.
- **Error Distribution:** The significant difference in MSE suggests that the ARIMA(1,1,2) model not only has smaller average errors but also fewer large errors compared to the ARIMA(0,1,8) model.

These results support the preference for the ARIMA(1,1,2) model due to its superior accuracy and robustness in predicting the data.

SECOND TIME SERIES

Now are going to load the second time series by using the “gtrends” package, and it is about the searches of the phrase “Chat GPT”.

Loading data and transformations

```
# Define the search term
search_term <- "Chat GPT"

# Calculate the current date and the date 5 years ago
end_date <- as.Date("2024-05-26")
start_date <- end_date %m-% years(5)

# Format the date range
time_frame <- paste0(start_date, " ", end_date)

# Fetch the trend data
trend_data <- gtrends(search_term, time = time_frame)

# Extract the interest over time data
interest_over_time <- trend_data$interest_over_time

gpt_data = interest_over_time[,1:2]


gpt_data$date <- as.Date(gpt_data$date, format = "%Y-%m-%d")

gpt_data$hits <- as.numeric(gsub("<1", "0", gpt_data$hits))
gpt_data$hits <- as.numeric(gpt_data$hits)

# Convert the data frame to xts object
gpt_xts <- xts(gpt_data$hits, order.by = gpt_data$date)
plot(gpt_xts,
     type = "l",
     col = "darkred",
     lwd = 3,
     main = "'Chat GPT' Search Trends - Last 5 years - Log")
```

'Chat GPT' Search Trends – Last 5 years – Log


2019-05-26 / 2024-05-26



```
# Log transformation
gpt_xts <- gpt_xts + 1e-8
log_gpt_xts <- log(gpt_xts)
```

Checking integration order

```
diff_log_gpt_xts <- diff(log_gpt_xts, differences = 1)
plot(diff_log_gpt_xts,
     type = "l",
     col = "darkred",
     lwd = 3,
     main = "'Chat GPT' Search Trends - Last 5 years - Log")
```



```
diff_log_gpt_xts <- diff_log_gpt_xts[-1,]
gpt_adf_test <- ur.df(diff_log_gpt_xts, type = "drift", lags = 0)
summary(gpt_adf_test)
```

```
#####
# Augmented Dickey-Fuller Test Unit Root Test #
#####
```

Test regression drift

Call:

```
lm(formula = z.diff ~ z.lag.1 + 1)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.4351	-0.0777	-0.0777	-0.0777	18.3430

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.07766	0.07129	1.089	0.277
z.lag.1	-0.87782	0.06179	-14.206	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.146 on 258 degrees of freedom

Multiple R-squared: 0.4389, Adjusted R-squared: 0.4367

F-statistic: 201.8 on 1 and 258 DF, p-value: < 2.2e-16

Value of test-statistic is: -14.2062 100.9084

Critical values for test statistics:

	1pct	5pct	10pct
tau2	-3.44	-2.87	-2.57
phi1	6.47	4.61	3.79

The test statistic is far below the critical value, so we can reject the null-hypothesis of non-stationarity.

Let's check more tests:

```
gpt_pp.test <- ur.pp(diff_log_gpt_xts, # tested series
                     type = c("Z-tau"), # standardization of the test
                     ↪ statistic needed
                     model = c("constant")) # constant deterministic
                     ↪ component
# which means we assume that any trends in the data are stochastic
summary(gpt_pp.test)
```

```
#####
# Phillips-Perron Unit Root Test #
#####
```

Test regression with intercept

Call:

```
lm(formula = y ~ y.l1)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.4351	-0.0777	-0.0777	-0.0777	18.3430

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.07766	0.07129	1.089	0.2770
y.l1	0.12218	0.06179	1.977	0.0491 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.146 on 258 degrees of freedom

Multiple R-squared: 0.01493, Adjusted R-squared: 0.01111

F-statistic: 3.91 on 1 and 258 DF, p-value: 0.04907

Value of test-statistic, type: Z-tau is: -14.2098

	aux. Z statistics
Z-tau-mu	1.0897

Critical values for Z statistics:

	1pct	5pct	10pct
critical values	-3.457006	-2.872754	-2.572697

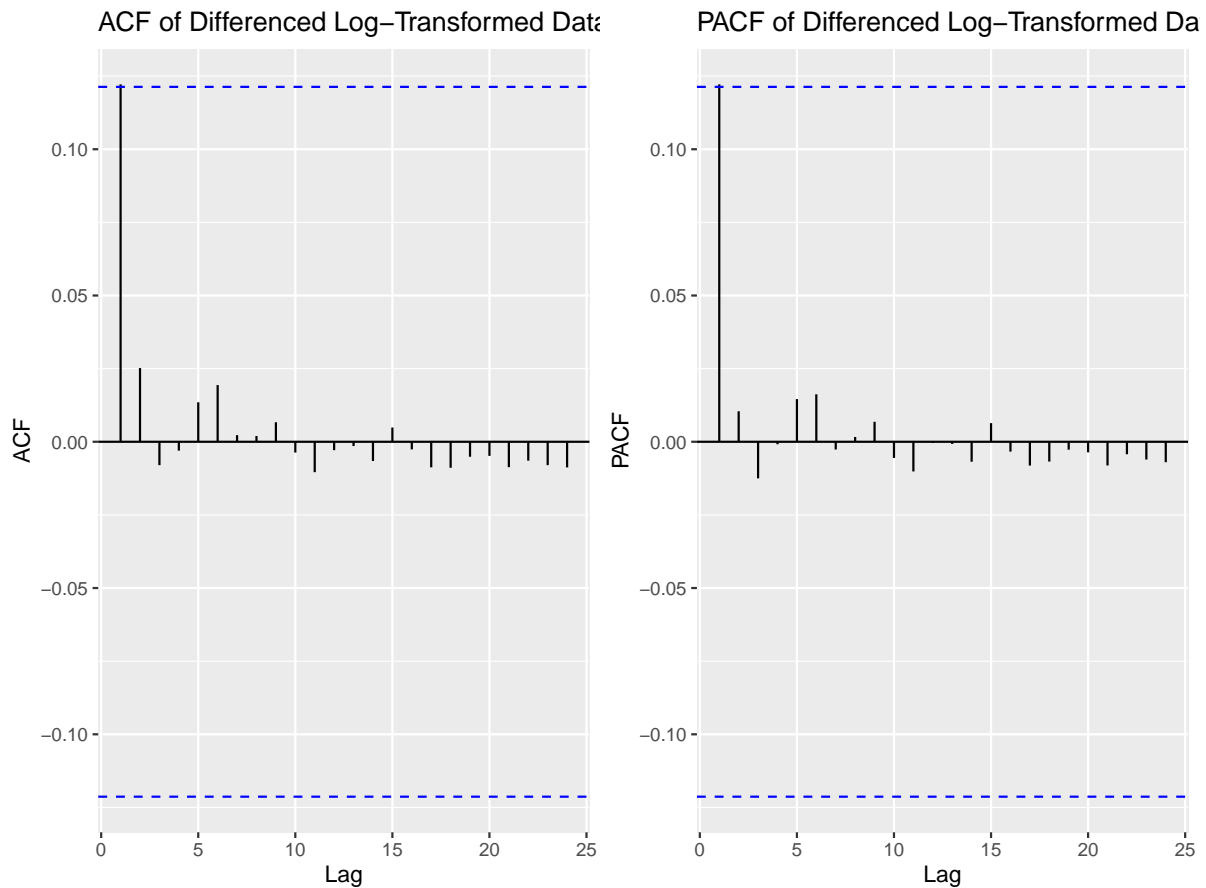
This test also confirms the non stationarity, so we can say that the series is $\sim I(1)$.

ACF - PACF Evaluations

```
# ACF and PACF plots of the differenced series
acf_plot <- ggAcf(diff_log_gpt_xts, main = "ACF of Differenced
  ↳ Log-Transformed Data")
pacf_plot <- ggPacf(diff_log_gpt_xts, main = "PACF of Differenced
  ↳ Log-Transformed Data")
```



```
gridExtra::grid.arrange(acf_plot, pacf_plot, ncol = 2)
```



From the plots we can see AR(1) and MA(1).

Let's now check ACF and PACF for the residuals:

ARIMA

Fitting two models

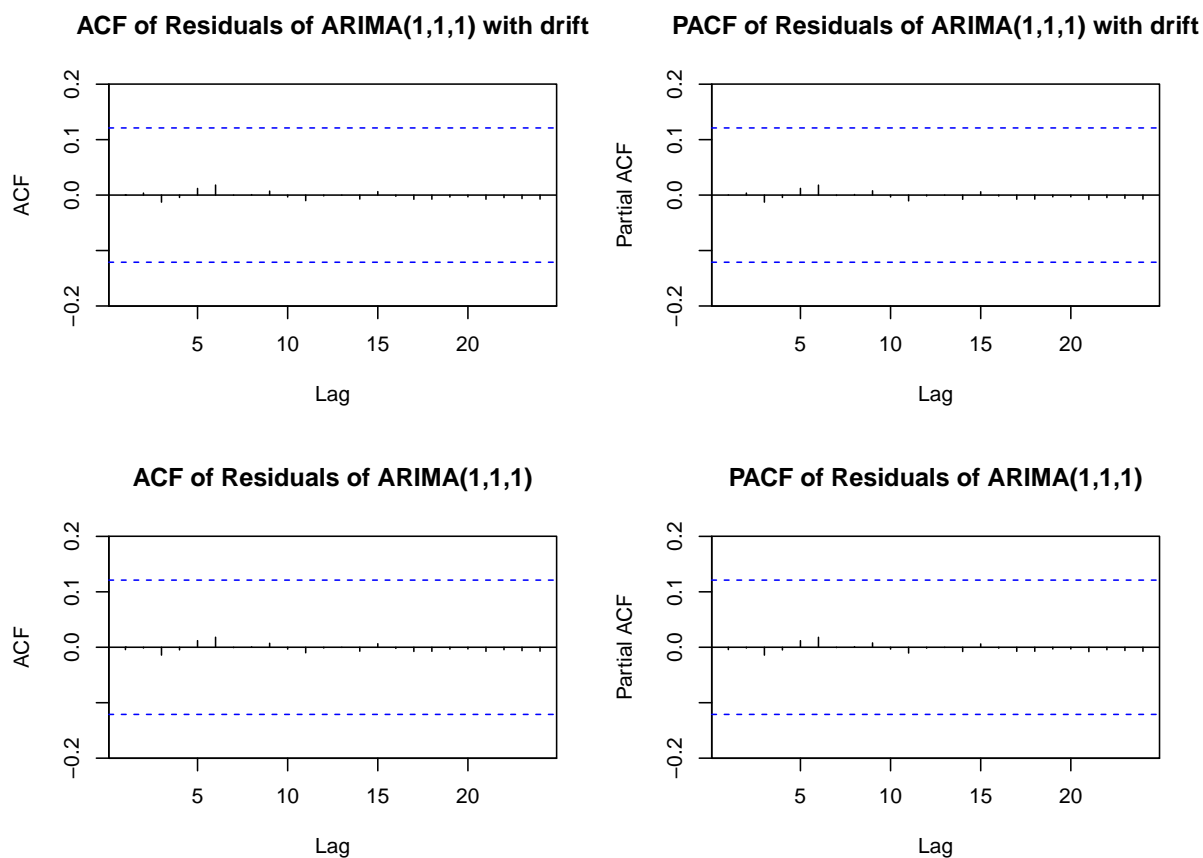
We are going to compare ARIMA (1,1,1) with drift and without it:

```
# Fit two different ARIMA models - constant and no constant
ts2_model1 <- Arima(log_gpt_xts, order = c(1, 1, 1), include.constant =
  ↪ TRUE)
ts2_model2 <- Arima(log_gpt_xts, order = c(1, 1, 1), include.constant =
  ↪ FALSE)
residuals_ts2_model1 <- residuals(ts2_model1)
residuals_ts2_model2 <- residuals(ts2_model2)
```

Diagnostics

Let's check their respective ACF and PACF plots:

```
par(mfrow = c(2, 2))
Acf(residuals_ts2_model1, main = "ACF of Residuals of ARIMA(1,1,1) with
  ↪ drift")
Pacf(residuals_ts2_model1, main = "PACF of Residuals of ARIMA(1,1,1)
  ↪ with drift")
Acf(residuals_ts2_model2, main = "ACF of Residuals of ARIMA(1,1,1)")
Pacf(residuals_ts2_model2, main = "PACF of Residuals of ARIMA(1,1,1)")
```



From the plots it looks like there is no autocorrelation in the residuals for both models, let's now check with the tests:

```
# Perform tests on residuals to check for white noise
lb_test_ts2_model1 <- Box.test(residuals_ts2_model1, lag = 1, type =
  ↪ "Ljung-Box")
lb_test_ts2_model2 <- Box.test(residuals_ts2_model2, lag = 1, type =
  ↪ "Ljung-Box")
bg_test_ts2_model1 <- bgtest(residuals_ts2_model1 ~ fitted(ts2_model1),
  ↪ order = 1)
bg_test_ts2_model2 <- bgtest(residuals_ts2_model2 ~ fitted(ts2_model2),
  ↪ order = 1)
```

After comparing lags = 5 (rule of thumb) and lags = 1 (intuition), 1 was selected since it makes more sense in the context of application.

```
# Create a table to compare results
comparison_table_2 <- data.frame(
  Model = c("ARIMA(1,1,1) with drift", "ARIMA(1,1,1)"),
  AIC = c(AIC(ts2_model1), AIC(ts2_model2)),
  BIC = c(BIC(ts2_model1), BIC(ts2_model2)),
  LjungBox_p_value = c(lb_test_ts2_model1$p.value,
    ↪ lb_test_ts2_model2$p.value), BG_statistic =
    ↪ c(bg_test_ts2_model1$p.value, bg_test_ts2_model2$p.value))

# Print the comparison table
print(comparison_table_2)
```

	Model	AIC	BIC	LjungBox_p_value	BG_statistic
1	ARIMA(1,1,1) with drift	816.8274	831.0855	0.9962941	0.9573842
2	ARIMA(1,1,1)	815.9925	826.6861	0.9460117	0.9858002

None of the models shows autocorrelation in the residuals, ARIMA (1,1,1) is selected for lower AIC and BIC values.

AUTO-ARIMA

Automatic model selection

Let's look for the best models in terms of AIC and BIC:

```
ts2_arma.best.AIC <-
  auto.arma(log_gpt_xts,
    d = 1, # parameter d of ARIMA model
    max.p = 2, # Maximum value of p
    max.q = 2, # Maximum value of q
    max.order = 4, # maximum p+q
    start.p = 0, # Starting value of p in stepwise
    ↪ procedure
    start.q = 0, # Starting value of q in stepwise
    ↪ procedure
    ic = "aic", # Information criterion to be used in
    ↪ model selection.
    stepwise = FALSE, # if FALSE considers all models
    allowdrift = TRUE, # include a constant
```

```
trace = TRUE)      # show summary of all models considered
```

Fitting models using approximations to speed things up...

```
ARIMA(0,1,0)           : 816.0132
ARIMA(0,1,0) with drift : 816.4792
ARIMA(0,1,1)           : 813.9386
ARIMA(0,1,1) with drift : 814.6977
ARIMA(0,1,2)           : 815.6699
ARIMA(0,1,2) with drift : 816.5003
ARIMA(1,1,0)           : 814.7477
ARIMA(1,1,0) with drift : 815.5497
ARIMA(1,1,1)           : 816.6958
ARIMA(1,1,1) with drift : 817.5255
ARIMA(1,1,2)           : 818.6703
ARIMA(1,1,2) with drift : 819.4913
ARIMA(2,1,0)           : 817.695
ARIMA(2,1,0) with drift : 818.5224
ARIMA(2,1,1)           : 819.6817
ARIMA(2,1,1) with drift : 820.5105
ARIMA(2,1,2)           : 821.6732
ARIMA(2,1,2) with drift : 822.4687
```

Now re-fitting the best model(s) without approximations...

Best model: ARIMA(0,1,1)

```
ts2_arma.best.BIC <-
  auto.arma(log_gpt_xts,
    d = 1,           # parameter d of ARIMA model
    max.p = 2,       # Maximum value of p
    max.q = 2,       # Maximum value of q
    max.order = 4,   # maximum p+q
    start.p = 0,     # Starting value of p in stepwise
    ↪ procedure
```

```

start.q = 0,          # Starting value of q in stepwise
  ↪ procedure
ic = "bic",           # Information criterion to be used in
  ↪ model selection.
stepwise = FALSE,     # if FALSE considers all models
allowdrift = TRUE,    # include a constant
trace = TRUE)         # show summary of all models considered

```

Fitting models using approximations to speed things up...

ARIMA(0,1,0)	: 819.5778
ARIMA(0,1,0) with drift	: 823.6083
ARIMA(0,1,1)	: 821.0676
ARIMA(0,1,1) with drift	: 825.3913
ARIMA(0,1,2)	: 826.3635
ARIMA(0,1,2) with drift	: 830.7584
ARIMA(1,1,0)	: 821.8767
ARIMA(1,1,0) with drift	: 826.2433
ARIMA(1,1,1)	: 827.3893
ARIMA(1,1,1) with drift	: 831.7836
ARIMA(1,1,2)	: 832.9284
ARIMA(1,1,2) with drift	: 837.3139
ARIMA(2,1,0)	: 828.3885
ARIMA(2,1,0) with drift	: 832.7805
ARIMA(2,1,1)	: 833.9398
ARIMA(2,1,1) with drift	: 838.3331
ARIMA(2,1,2)	: 839.4958
ARIMA(2,1,2) with drift	: 843.8558

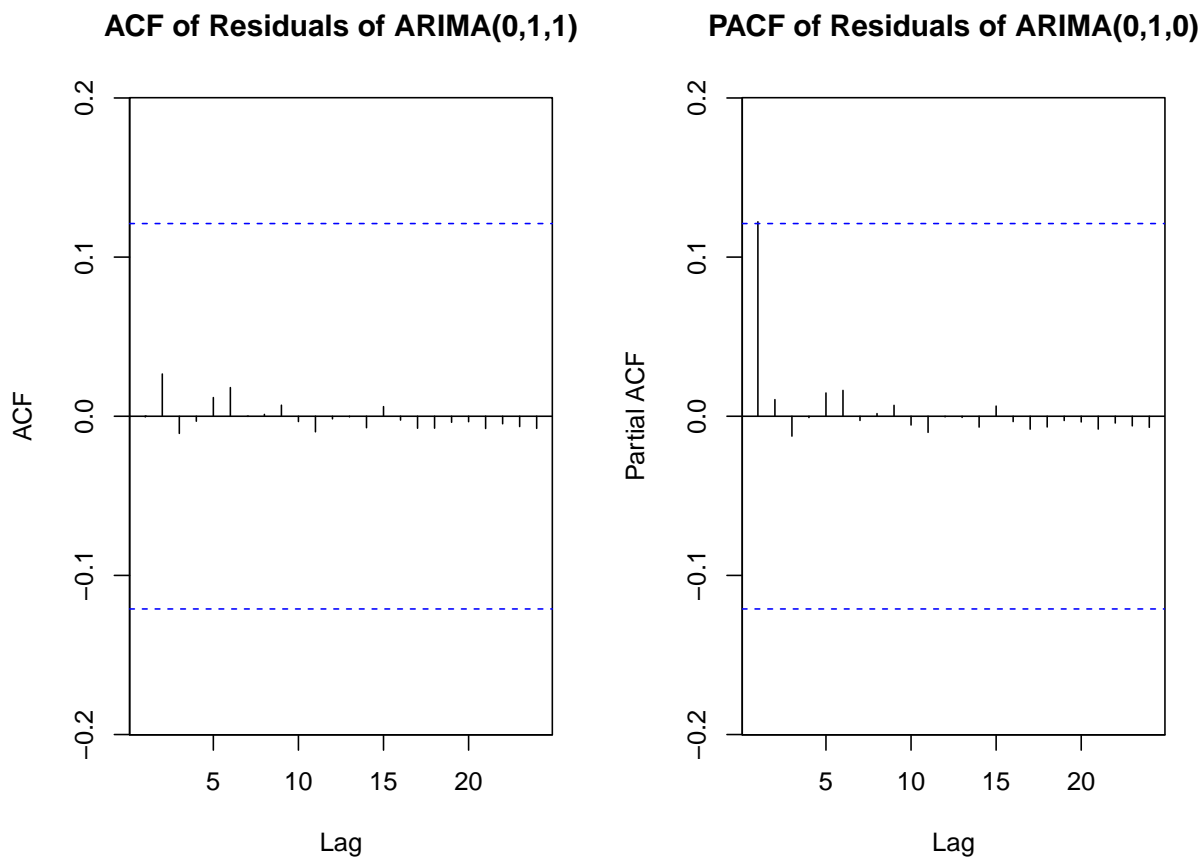
Now re-fitting the best model(s) without approximations...

Best model: ARIMA(0,1,0)

From these automatic selection methods, ARIMA (0,1,1) and ARIMA (0,1,0) were selected.

Automatic models diagnostics

```
ts2_arma.best.AIC_res = residuals(ts2_arma.best.AIC)
ts2_arma.best.BIC_res = residuals(ts2_arma.best.BIC)
par(mfrow = c(1, 2))
Acf(ts2_arma.best.AIC_res, main = "ACF of Residuals of ARIMA(0,1,1)")
Pacf(ts2_arma.best.BIC_res, main = "PACF of Residuals of ARIMA(0,1,0)")
```



ARIMA (0,1,1) shows no autocorrelation in the residuals while ARIMA (0,1,0) shows just one spike at lag 1.

Let's continue with the tests:

```
lb_test_ts2_model3 <- Box.test(ts2_arma.best.AIC_res, lag = 1, type =  
  ↪ "Ljung-Box")
```

```

lb_test_ts2_model4 <- Box.test(ts2_arma.best.BIC_res, lag = 1, type =
  ↪ "Ljung-Box")
bg_test_ts2_model3 <- bgtest(ts2_arma.best.AIC_res ~
  ↪ fitted(ts2_arma.best.AIC), order = 1)
bg_test_ts2_model4 <- bgtest(ts2_arma.best.BIC_res ~
  ↪ fitted(ts2_arma.best.BIC), order = 1)

# Create a table to compare results
comparison_table_3 <- data.frame(
  Model = c("ARIMA(0,1,1)", "ARIMA(0,1,0)"),
  AIC = c(AIC(ts2_arma.best.AIC), AIC(ts2_arma.best.BIC)),
  BIC = c(BIC(ts2_arma.best.AIC), BIC(ts2_arma.best.BIC)),
  LjungBox_p_value = c(lb_test_ts2_model3$p.value,
    ↪ lb_test_ts2_model4$p.value), BG_statistic =
    ↪ c(bg_test_ts2_model3$p.value, bg_test_ts2_model4$p.value))

# Print the comparison table
print(comparison_table_3)

```

	Model	AIC	BIC	LjungBox_p_value	BG_statistic
1	ARIMA(0,1,1)	814.2354	821.3645	0.99479494	0.96827473
2	ARIMA(0,1,0)	816.2953	819.8598	0.04666743	0.04370179

The tests confirm the presence of autocorrelation for the ARIMA (0,1,0) when considering one lag.

For this reason, after checking for the significance of its coefficient we are going to select ARIMA(0,1,1)

```
coeftest(ts2_arma.best.AIC)
```

z test of coefficients:

```

      Estimate Std. Error z value Pr(>|z|)
ma1 0.121174    0.059452  2.0382  0.04153 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

As we can see the moving average term is significant at 5% level.

FORECASTING

Manual selection

Let's now compare ARIMA(1,1,1) and ARIMA(0,1,1) models in terms of forecasting:

```
# Split the data into training and testing sets
split_point <- floor(0.95 * length(log_gpt_xts))
ts2_train_data <- window(log_gpt_xts, end =
  ↪ index(log_gpt_xts)[split_point])
ts2_test_data <- window(log_gpt_xts, start =
  ↪ index(log_gpt_xts)[split_point + 1])

ts2_manual_model <- Arima(ts2_train_data, order = c(1, 1, 1))
ts2_auto_model <- Arima(ts2_train_data, order = c(0,1,1))


# Generate out-of-sample forecasts for the length of the test set
forecast_horizon <- length(ts2_test_data)
ts2_out_of_sample_forecast <- forecast(ts2_manual_model, h =
  ↪ forecast_horizon)
ts2_forecasts_data <- data.frame(f_mean =
  ↪ as.numeric(ts2_out_of_sample_forecast$mean),
                                f_lower =
  ↪ as.numeric(ts2_out_of_sample_forecast$lower[,
  ↪ 2]),
                                f_upper =
  ↪ as.numeric(ts2_out_of_sample_forecast$upper[,
  ↪ 2]))

gpt.oos <- ts2_test_data
# Ensure names match for consistency
names(gpt.oos) <- "V1"
colnames(ts2_train_data) <- "V1"
gpt2 <- rbind(ts2_train_data[, "V1"], gpt.oos)
ts2_forecasts_xts <- xts(ts2_forecasts_data, order.by = index(gpt.oos))
ts2_forecasted_data <- merge(gpt2, ts2_forecasts_xts)
```

```
plot(ts2_forecasted_data[, c("V1", "f_mean", "f_lower", "f_upper")],
     major.ticks = "years",
     grid.ticks.on = "years",
     grid.ticks.lty = 3,
     main = "forecast of 'Chat GPT' search term",
     col = c("black", "blue", "red", "red"))
```

forecast of 'Chat GPT' search term

2019-05-26 / 2024-05-26



```
eval_data_oos <- tail(ts2_forecasted_data, length(ts2_test_data))
eval_data_oos$V1 <- exp(eval_data_oos$V1)
eval_data_oos$f_mean <- exp(eval_data_oos$f_mean)
eval_data_oos$f_lower <- exp(eval_data_oos$f_lower)
eval_data_oos$f_upper <- exp(eval_data_oos$f_upper)
eval_data_oos$mae <- abs(eval_data_oos$V1 - eval_data_oos$f_mean)
```

```
eval_data_oos$mse <- (eval_data_oos$V1 - eval_data_oos$f_mean) ^ 2
eval_data_oos$mape <- abs((eval_data_oos$V1 -
  ↪ eval_data_oos$f_mean)/eval_data_oos$V1)
eval_data_oos$amape <- abs((eval_data_oos$V1 -
  ↪ eval_data_oos$f_mean)/(eval_data_oos$V1 + eval_data_oos$f_mean))
eval_data_oos
```

	V1	f_mean	f_lower	f_upper	mae	mse
2024-02-25	71	72.10376	7.157822941	726.3315	1.1037582	1.2182821
2024-03-03	75	72.35574	2.234929701	2342.5139	2.6442550	6.9920846
2024-03-10	75	72.41285	0.906070756	5787.2097	2.5871457	6.6933228
2024-03-17	77	72.42578	0.426150767	12309.0076	4.5742240	20.9235255
2024-03-24	66	72.42870	0.220575112	23782.9025	6.4286985	41.3281649
2024-03-31	72	72.42936	0.122108159	42962.0114	0.4293595	0.1843496
2024-04-07	76	72.42951	0.071079923	73804.7194	3.5704910	12.7484061
2024-04-14	80	72.42954	0.043034354	121903.5060	7.5704572	57.3118224
2024-04-21	83	72.42955	0.026896776	195043.4445	10.5704496	111.7344041
2024-04-28	82	72.42955	0.017261030	303923.9292	9.5704478	91.5934718
2024-05-05	83	72.42955	0.011328466	463084.7773	10.5704474	111.7343592
2024-05-12	94	72.42955	0.007580028	692087.1290	21.5704474	465.2841993
2024-05-19	100	72.42955	0.005158294	1017010.6833	27.5704473	760.1295665
2024-05-26	98	72.42955	0.003563058	1472342.0819	25.5704473	653.8477769
	mape	amape				
2024-02-25	0.015545890	0.007712992				
2024-03-03	0.035256734	0.017944703				
2024-03-10	0.034495276	0.017550340				
2024-03-17	0.059405507	0.030612015				
2024-03-24	0.097404523	0.046440504				
2024-03-31	0.005963326	0.002972799				
2024-04-07	0.046980145	0.024055129				
2024-04-14	0.094630715	0.049665288				
2024-04-21	0.127354814	0.068007979				
2024-04-28	0.116712779	0.061972904				
2024-05-05	0.127354789	0.068007964				
2024-05-12	0.229472844	0.129607074				
2024-05-19	0.275704473	0.159893979				
2024-05-26	0.260922932	0.150035290				


```
ts2_perf_manual_model = colMeans(eval_data_oos[, c("mae", "mse", "mape",
  ↪ "amape")])
```

Automatic selection

```
# Generate out-of-sample forecasts for the length of the test set
forecast_horizon <- length(ts2_test_data)
ts2_out_of_sample_forecast <- forecast(ts2_auto_model, h =
  ↪ forecast_horizon)
ts2_forecasts_data_auto <- data.frame(f_mean =
  ↪ as.numeric(ts2_out_of_sample_forecast$mean),
                                     f_lower =
  ↪ as.numeric(ts2_out_of_sample_forecast$lower[,
  ↪ 2]),
                                     f_upper =
  ↪ as.numeric(ts2_out_of_sample_forecast$upper[,
  ↪ 2]))

ts2_forecasts_auto_xts <- xts(ts2_forecasts_data_auto,
  order.by = index(gpt.oos))
ts2_forecasted_data_auto <- merge(gpt2, ts2_forecasts_auto_xts)

plot(ts2_forecasted_data_auto[, c("V1", "f_mean", "f_lower",
  ↪ "f_upper")],
     major.ticks = "years",
     grid.ticks.on = "years",
     grid.ticks.lty = 3,
     main = "forecast of 'Chat GPT' search term",
     col = c("black", "blue", "red", "red"))
```



```
eval_data_oos <- tail(ts2_forecasted_data_auto, length(ts2_test_data))
eval_data_oos$V1 <- exp(eval_data_oos$V1)
eval_data_oos$f_mean <- exp(eval_data_oos$f_mean)
eval_data_oos$f_lower <- exp(eval_data_oos$f_lower)
eval_data_oos$f_upper <- exp(eval_data_oos$f_upper)
eval_data_oos$mae <- abs(eval_data_oos$V1 - eval_data_oos$f_mean)
eval_data_oos$mse <- (eval_data_oos$V1 - eval_data_oos$f_mean) ^ 2
eval_data_oos$mape <- abs((eval_data_oos$V1 -
  ↪ eval_data_oos$f_mean)/eval_data_oos$V1)
eval_data_oos$amape <- abs((eval_data_oos$V1 -
  ↪ eval_data_oos$f_mean)/(eval_data_oos$V1 + eval_data_oos$f_mean))
eval_data_oos
```

	V1	f_mean	f_lower	f_upper	mae	mse
2024-02-25	71	72.00684	7.174120497	722.7345	1.006836196	1.013719e+00

2024-03-03	75	72.00684	2.252243337	2302.1422	2.993163804	8.959030e+00			
2024-03-10	75	72.00684	0.954494376	5432.1792	2.993163804	8.959030e+00			
2024-03-17	77	72.00684	0.467276408	11096.1828	4.993163804	2.493168e+01			
2024-03-24	66	72.00684	0.250147593	20727.7008	6.006836196	3.608208e+01			
2024-03-31	72	72.00684	0.142535920	36376.6863	0.006836196	4.673358e-05			
2024-04-07	76	72.00684	0.085106877	60923.2135	3.993163804	1.594536e+01			
2024-04-14	80	72.00684	0.052718327	98352.5980	7.993163804	6.389067e+01			
2024-04-21	83	72.00684	0.033645282	154107.3280	10.993163804	1.208497e+02			
2024-04-28	82	72.00684	0.022013706	235534.3726	9.993163804	9.986332e+01			
2024-05-05	83	72.00684	0.014711248	352450.3364	10.993163804	1.208497e+02			
2024-05-12	94	72.00684	0.010012491	517851.5895	21.993163804	4.836993e+02			
2024-05-19	100	72.00684	0.006924373	748802.0354	27.993163804	7.836172e+02			
2024-05-26	98	72.00684	0.004856957	1067537.7247	25.993163804	6.756446e+02			
		mape	amape						
2024-02-25	1.418079e-02	7.040476e-03							
2024-03-03	3.990885e-02	2.036071e-02							
2024-03-10	3.990885e-02	2.036071e-02							
2024-03-17	6.484628e-02	3.350963e-02							
2024-03-24	9.101267e-02	4.352564e-02							
2024-03-31	9.494717e-05	4.747133e-05							
2024-04-07	5.254163e-02	2.697959e-02							
2024-04-14	9.991455e-02	5.258424e-02							
2024-04-21	1.324478e-01	7.092051e-02							
2024-04-28	1.218679e-01	6.488779e-02							
2024-05-05	1.324478e-01	7.092051e-02							
2024-05-12	2.339698e-01	1.324835e-01							
2024-05-19	2.799316e-01	1.627445e-01							
2024-05-26	2.652364e-01	1.528948e-01							

```
ts2_perf_auto_model = colMeans(eval_data_oos[, c("mae", "mse", "mape",
↪ "amape")])
```

Comparison

```
# Combine the performance metrics into a single data frame
ts2_combined_perf <- data.frame(
  ARIMA_1_1_1 = ts2_perf_manual_model,
  ARIMA_0_1_1 = ts2_perf_auto_model
)
```

```
# Print the combined table
print(ts2_combined_perf)
```

	ARIMA_1_1_1	ARIMA_0_1_1
mae	9.59507686	9.85323646
mse	167.26598112	174.59323412
mape	0.10908605	0.11202213
amape	0.05960564	0.06137572

As we can see from the table, it looks like ARIMA (1,1,1) performs slightly better, so also in this case the manual model proved to be better.

TESTING COINTEGRATION


First of all, we are going to shorten the period considered up to the date when ChatGPT went out.

```
series <- merge.xts(log_trend1_xts, log_gpt_xts)
colnames(series) <- c("AI", "ChatGPT")
```

```
plot(series[184:262, 1:2],
      col = c("black", "blue"),
      major.ticks = "months",
      grid.ticks.on = "months",
      grid.ticks.lty = 7,
      main = "'AI' vs 'ChatGPT' search trends",
      legend.loc = "bottomright")
```

'AI' vs 'ChatGPT' search trends

2022-11-27 / 2024-05-26



The cointegrating vector is then estimated:

```
model.coint <- lm(AI ~ ChatGPT, data = series)
summary(model.coint)
```

Call:

```
lm(formula = AI ~ ChatGPT, data = series)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.79375	-0.10428	0.00319	0.14704	0.39216

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.94165	0.02053	191.98	<2e-16 ***
ChatGPT	0.06505	0.00132	49.27	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2196 on 260 degrees of freedom

Multiple R-squared: 0.9033, Adjusted R-squared: 0.9029

F-statistic: 2428 on 1 and 260 DF, p-value: < 2.2e-16

Testing the stationarity of the residuals:

```
adf.test(residuals(model.coint))
```

Augmented Dickey-Fuller Test

```
data: residuals(model.coint)
```

```
Dickey-Fuller = -3.5227, Lag order = 6, p-value = 0.04104
```

```
alternative hypothesis: stationary
```

Since the p-value is 0.04, we reject the null about non-stationarity, so we can conclude that the two series are **cointegrated**.

Johansen cointegration test

We can also check it by performing the Johansen Test:

```
# Convert the relevant columns to a matrix
series_matrix <- as.matrix(series[, c("ChatGPT", "AI")])

# Perform the Johansen cointegration test
johansen_test <- ca.jo(series_matrix, type = "trace", ecdet = "none", K
↪ = 5)

# Summary of the test results
summary(johansen_test)
```

```
#####
# Johansen-Procedure #
#####
```

Test type: trace statistic , with linear trend

```
Eigenvalues (lambda):
[1] 0.0874720246 0.0007934656
```

Values of teststatistic and critical values of test:

```

          test 10pct  5pct  1pct
r <= 1 |   0.20   6.50   8.18 11.65
r = 0  |  23.73  15.66  17.95 23.52
```

Eigenvectors, normalised to first column:
(These are the cointegration relations)

```

          ChatGPT.15      AI.15
ChatGPT.15   1.00000   1.00000
AI.15        -12.39133 -26.45684
```

Weights W:
(This is the loading matrix)

```

          ChatGPT.15      AI.15
ChatGPT.d -0.040607903 -3.328168e-03
AI.d       0.004530529 -8.673299e-05
```

Null Hypothesis $r=0$:

- The test statistic is 23.55, which is above the 5% critical value of 14.90. This suggests that we can reject the null hypothesis of no cointegrating relationships at the 5% level, indicating the presence of one cointegrating relationship between ChatGPT and AI.

Null Hypothesis $r \leq 1$:

- The test statistic is 0.20, which is well below the 5% critical value of 8.18. This means we fail to reject the null hypothesis that there is at most 1 cointegrating relationship.

We can conclude that the two time series are cointegrated and there is only one cointegrating vector.

GRANGER CAUSALITY

In this section we want to check if either of the series is Granger caused by the other one:

```
diff_series <- diff(series)

granger_test_result <- grangertest(AI ~ ChatGPT, order = 5, data=
  ↪ diff_series)
print(granger_test_result)
```

Granger causality test

Model 1: AI ~ Lags(AI, 1:5) + Lags(ChatGPT, 1:5)

Model 2: AI ~ Lags(AI, 1:5)

	Res.Df	Df	F	Pr(>F)
1	245			
2	250	-5	11.825	2.943e-10 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The output of the Granger causality test shows that the p-value is close to 0, which is lower than the typical significance level of 0.05. This indicates that we reject the null hypothesis of no granger causality, suggesting that **AI** search **does** Granger-cause **Chat GPT** searches.

In other words, including lagged values of AI searches time series improves the predictive ability of the model over including only the ones of ChatGPT searches time series.

```
granger_test_result_reverse <- grangertest(ChatGPT ~ AI, order = 5, data
  ↪ = diff_series)
print(granger_test_result_reverse)
```

Granger causality test

Model 1: ChatGPT ~ Lags(ChatGPT, 1:5) + Lags(AI, 1:5)

Model 2: ChatGPT ~ Lags(ChatGPT, 1:5)

	Res.Df	Df	F	Pr(>F)
1	245			
2	250	-5	0.2445	0.9423

Given the p-value higher than 0.05, we fail to reject the null hypothesis, meaning that the past values of **Chat GPT** searches **do not** Granger-cause **AI** searches.

VAR MODEL

Selecting lags

First of all, the optimal number of lags has to be found:

```
VARselect(series, # input data for VAR
          lag.max = 6)      # maximum lag
```

\$selection

AIC(n)	HQ(n)	SC(n)	FPE(n)
5	3	2	5

\$criteria

	1	2	3	4	5
AIC(n)	-5.283402284	-5.417711694	-5.443040351	-5.449677753	-5.48101290
HQ(n)	-5.249983755	-5.362014146	-5.365063784	-5.34942194	-5.35847830
SC(n)	-5.200312188	-5.279228200	-5.249163460	-5.20040724	-5.17634922
FPE(n)	0.005075145	0.004437333	0.004326428	0.00429794	0.00416555
	6				
AIC(n)	-5.455769282				
HQ(n)	-5.310955656				
SC(n)	-5.095712198				
FPE(n)	0.004272337				

```
VARselect(series, lag.max = 6) %>%
  .$criteria %>%
  t() %>%
  as_tibble() %>%
  mutate(nLags = 1:nrow(.)) %>%
  select(nLags, everything()) %>%
  kbl(digits = 3) %>%
  kable_classic("striped", full_width = F)
```

Since 2,3 and 5 lags are optimal, 5 are going to be selected.

nLags	AIC(n)	HQ(n)	SC(n)	FPE(n)
1	-5.283	-5.250	-5.200	0.005
2	-5.418	-5.362	-5.279	0.004
3	-5.443	-5.365	-5.249	0.004
4	-5.450	-5.349	-5.200	0.004
5	-5.481	-5.358	-5.176	0.004
6	-5.456	-5.311	-5.096	0.004

Model

```
series.var5 <- VAR(series, p = 5) # order of VAR model
summary(series.var5)
```

VAR Estimation Results:

=====

Endogenous variables: AI, ChatGPT

Deterministic variables: const

Sample size: 257

Log Likelihood: -1.956

Roots of the characteristic polynomial:

1.003 0.9109 0.6321 0.6321 0.573 0.573 0.5611 0.5611 0.4603 0.4603

Call:

VAR(y = series, p = 5)

Estimation results for equation AI:

=====

AI = AI.l1 + ChatGPT.l1 + AI.l2 + ChatGPT.l2 + AI.l3 + ChatGPT.l3 + AI.l4 + ChatGPT.l4 + AI.l5 + ChatGPT.l5

	Estimate	Std. Error	t value	Pr(> t)
AI.l1	0.661396	0.063066	10.487	< 2e-16 ***
ChatGPT.l1	0.022964	0.002998	7.661	4.25e-13 ***
AI.l2	0.106046	0.075941	1.396	0.163849
ChatGPT.l2	-0.018539	0.004497	-4.123	5.12e-05 ***
AI.l3	0.233649	0.074797	3.124	0.002000 **
ChatGPT.l3	-0.008931	0.004652	-1.920	0.056001 .
AI.l4	-0.107777	0.072572	-1.485	0.138796
ChatGPT.l4	-0.001905	0.004668	-0.408	0.683505
AI.l5	0.052842	0.058258	0.907	0.365276
ChatGPT.l5	0.010856	0.003307	3.283	0.001176 **

```
const      0.234435    0.066732    3.513 0.000527 ***
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 0.05395 on 246 degrees of freedom

Multiple R-Squared: 0.9944, Adjusted R-squared: 0.9942

F-statistic: 4376 on 10 and 246 DF, p-value: < 2.2e-16

Estimation results for equation ChatGPT:

```
=====
```

ChatGPT = AI.l1 + ChatGPT.l1 + AI.l2 + ChatGPT.l2 + AI.l3 + ChatGPT.l3 + AI.l4 + ChatGPT.l4 +

	Estimate	Std. Error	t value	Pr(> t)
AI.l1	1.320837	1.360791	0.971	0.333
ChatGPT.l1	1.100036	0.064682	17.007	<2e-16 ***
AI.l2	-0.722181	1.638603	-0.441	0.660
ChatGPT.l2	-0.126916	0.097029	-1.308	0.192
AI.l3	0.045849	1.613917	0.028	0.977
ChatGPT.l3	-0.008105	0.100367	-0.081	0.936
AI.l4	0.812773	1.565899	0.519	0.604
ChatGPT.l4	0.017738	0.100720	0.176	0.860
AI.l5	-0.866039	1.257050	-0.689	0.492
ChatGPT.l5	-0.026689	0.071348	-0.374	0.709
const	-2.334692	1.439896	-1.621	0.106

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 1.164 on 246 degrees of freedom

Multiple R-Squared: 0.9878, Adjusted R-squared: 0.9874

F-statistic: 1999 on 10 and 246 DF, p-value: < 2.2e-16

Covariance matrix of residuals:

```
      AI ChatGPT
AI      0.00291 0.01204
ChatGPT 0.01204 1.35504
```

Correlation matrix of residuals:

```
      AI ChatGPT
```

AI	1.0000	0.1917
ChatGPT	0.1917	1.0000

The roots of the characteristic polynomial are within or very close to the unit circle (1.002, 0.9122, etc.), indicating that the VAR model is stable.

Equation for AI:

- Significant lags: **AI.l1**, **ChatGPT.l1**, **ChatGPT.l2**, **AI.l3**, **ChatGPT.l5**, and **const**, indicating a good predictive power of ChatGPT on AI.
- High R^2 (0.9944) indicates that the model explains 99.44% of the variance in AI.

Equation for ChatGPT:

- Significant lags: only **ChatGPT.l1**, while the other being non significant; this indicates limited predictive power of AI on ChatGPT.

Residuals:

- The residuals show some correlation (0.1912), but this is not particularly strong.

Diagnostics

```
plot(series.var5)
```

Diagram of fit and residuals for AI


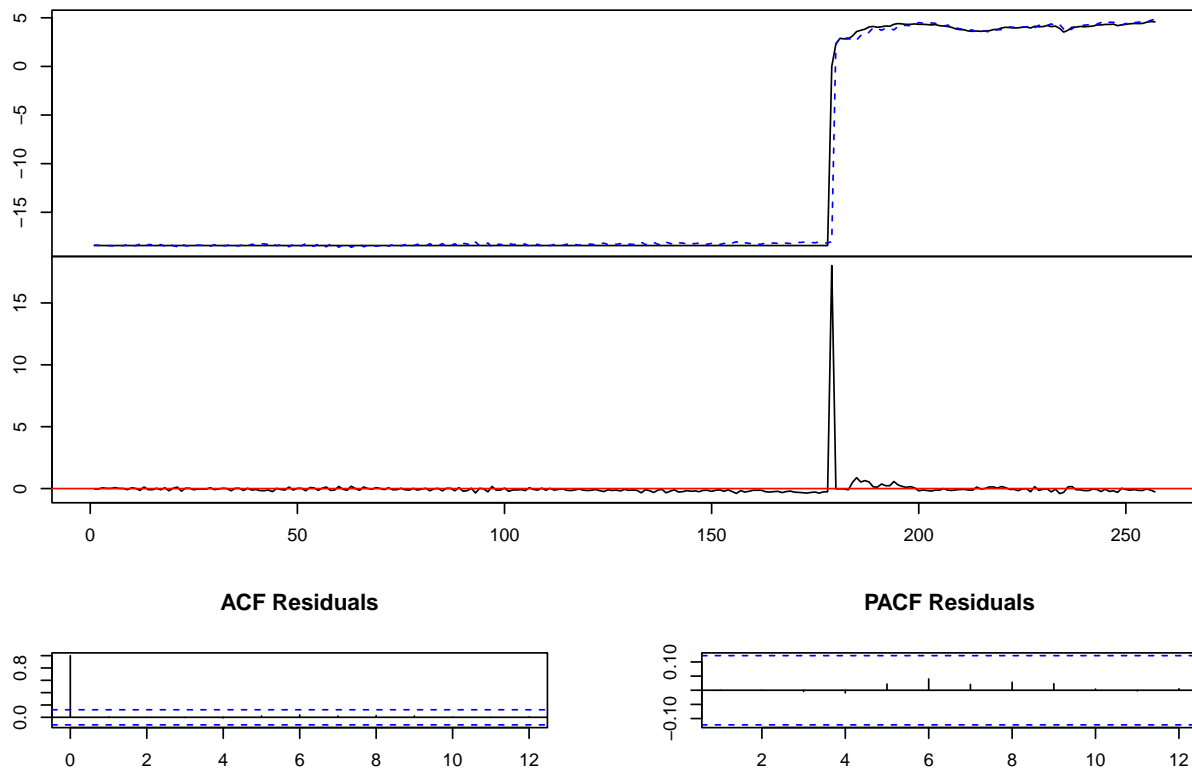


Diagram of fit and residuals for ChatGPT



For both equations we see significant spikes at lag 0 of the ACF plot. This autocorrelation at lag 0 indicates that the current values of each variable strongly depend on their previous values, which can happen in time series data where variables exhibit persistence.

Let's check the residuals formally by using **Portmanteau test**:

```
serial.test(series.var5)
```

Portmanteau Test (asymptotic)

```
data: Residuals of VAR object series.var5  
Chi-squared = 36.008, df = 44, p-value = 0.7988
```

Since the P-value is 0.77, we fail to reject the null of no autocorrelation.

Also with the BG test:

```
serial.test(series.var5, type = "BG")
```

Breusch-Godfrey LM test

```
data: Residuals of VAR object series.var5  
Chi-squared = 13.791, df = 20, p-value = 0.8409
```

Also in this case we fail to reject the null hypothesis since the p-value is > 0.05 .

Information Criteria

Let's compare with the information criteria with models having the other suggested number of lags (2 and 3):

```
series.var2 = VAR(series, p=2)  
series.var3 = VAR(series, p=3)
```

```
AIC(series.var2, series.var3, series.var5)
```

	df	AIC
series.var2	10	58.78234
series.var3	14	53.97275
series.var5	22	47.91246

```
BIC(series.var2, series.var3, series.var5)
```


	df	BIC
series.var2	10	94.38916
series.var3	14	103.76834
series.var5	22	125.99213

AIC suggests 5 or 3 lags, BIC suggests 2 or 3 lags.


Since BIC is slightly preferred in the context of VAR, let's check the residuals of the model with 2 lags:

```
plot(series.var2)
```

Diagram of fit and residuals for AI



ACF Residuals



PACF Residuals


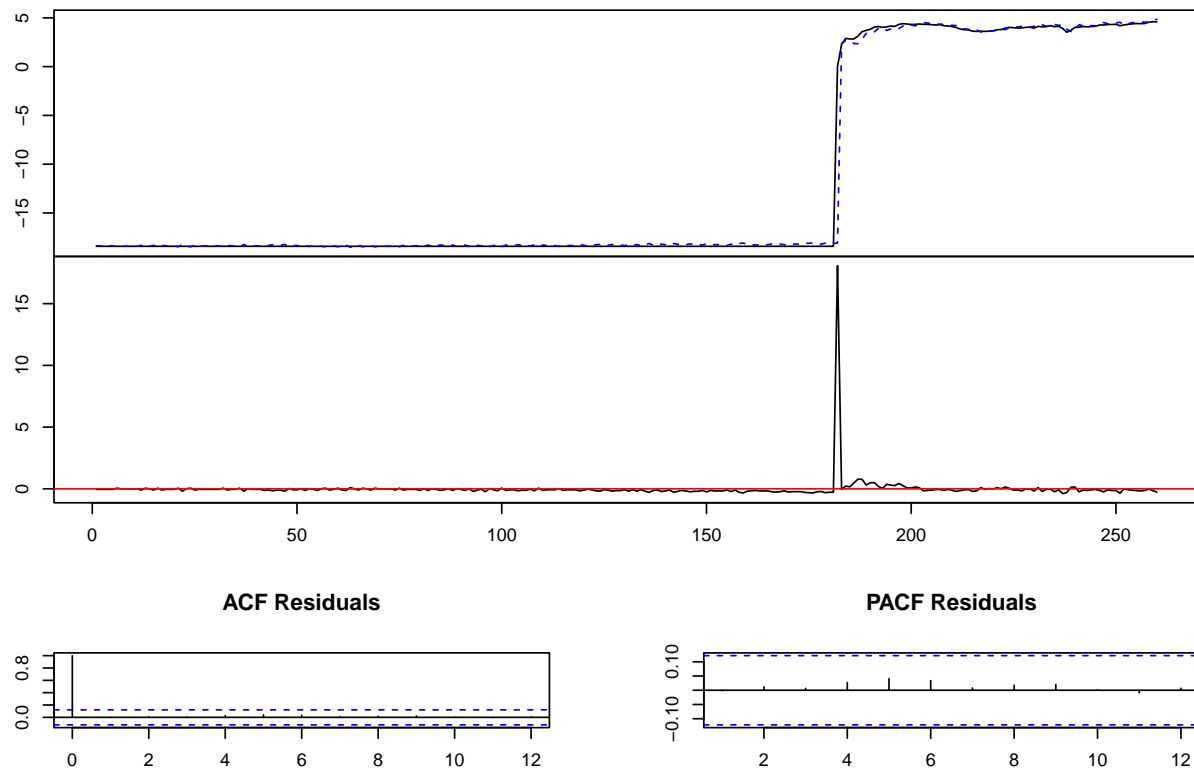


Diagram of fit and residuals for ChatGPT



A spike at lag 2 of the PACF plot can be seen; let's test it formally:

```
serial.test(series.var2)
```

Portmanteau Test (asymptotic)

data: Residuals of VAR object series.var2
Chi-squared = 63.769, df = 56, p-value = 0.2221

```
serial.test(series.var2, type = "BG")
```

Breusch-Godfrey LM test

```
data: Residuals of VAR object series.var2
Chi-squared = 40.967, df = 20, p-value = 0.003762
```

As we can see from the Breusch-Godfrey test, in this case the residuals show autocorrelation, so the model with 5 lags is still preferred.

The VAR model with 3 lags is also checked:

```
summary(series.var3)
```

VAR Estimation Results:

=====

Endogenous variables: AI, ChatGPT

Deterministic variables: const

Sample size: 259

Log Likelihood: -12.986

Roots of the characteristic polynomial:

1.003 0.9384 0.4494 0.4494 0.2324 0.09186

Call:

VAR(y = series, p = 3)

Estimation results for equation AI:

=====

AI = AI.l1 + ChatGPT.l1 + AI.l2 + ChatGPT.l2 + AI.l3 + ChatGPT.l3 + const

	Estimate	Std. Error	t value	Pr(> t)
AI.l1	0.7127668	0.0621562	11.467	< 2e-16 ***
ChatGPT.l1	0.0222271	0.0030903	7.193	7.24e-12 ***
AI.l2	0.0433975	0.0740655	0.586	0.558444
ChatGPT.l2	-0.0200309	0.0046093	-4.346	2.01e-05 ***
AI.l3	0.2126966	0.0586075	3.629	0.000344 ***
ChatGPT.l3	0.0005411	0.0033523	0.161	0.871892
const	0.1404485	0.0657533	2.136	0.033644 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.05573 on 252 degrees of freedom

Multiple R-Squared: 0.9939, Adjusted R-squared: 0.9938

F-statistic: 6863 on 6 and 252 DF, p-value: < 2.2e-16

Estimation results for equation ChatGPT:

=====

ChatGPT = AI.l1 + ChatGPT.l1 + AI.l2 + ChatGPT.l2 + AI.l3 + ChatGPT.l3 + const

	Estimate	Std. Error	t value	Pr(> t)
AI.l1	1.18011	1.28485	0.918	0.359
ChatGPT.l1	1.10196	0.06388	17.250	<2e-16 ***
AI.l2	-0.77895	1.53103	-0.509	0.611
ChatGPT.l2	-0.12215	0.09528	-1.282	0.201
AI.l3	0.13245	1.21149	0.109	0.913
ChatGPT.l3	-0.01875	0.06930	-0.271	0.787
const	-2.08162	1.35920	-1.532	0.127

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.152 on 252 degrees of freedom

Multiple R-Squared: 0.9878, Adjusted R-squared: 0.9876

F-statistic: 3414 on 6 and 252 DF, p-value: < 2.2e-16

Covariance matrix of residuals:

	AI	ChatGPT
AI	0.003106	0.01093
ChatGPT	0.010927	1.32724

Correlation matrix of residuals:

	AI	ChatGPT
AI	1.0000	0.1702
ChatGPT	0.1702	1.0000

```
plot(series.var3)
```

Diagram of fit and residuals for AI

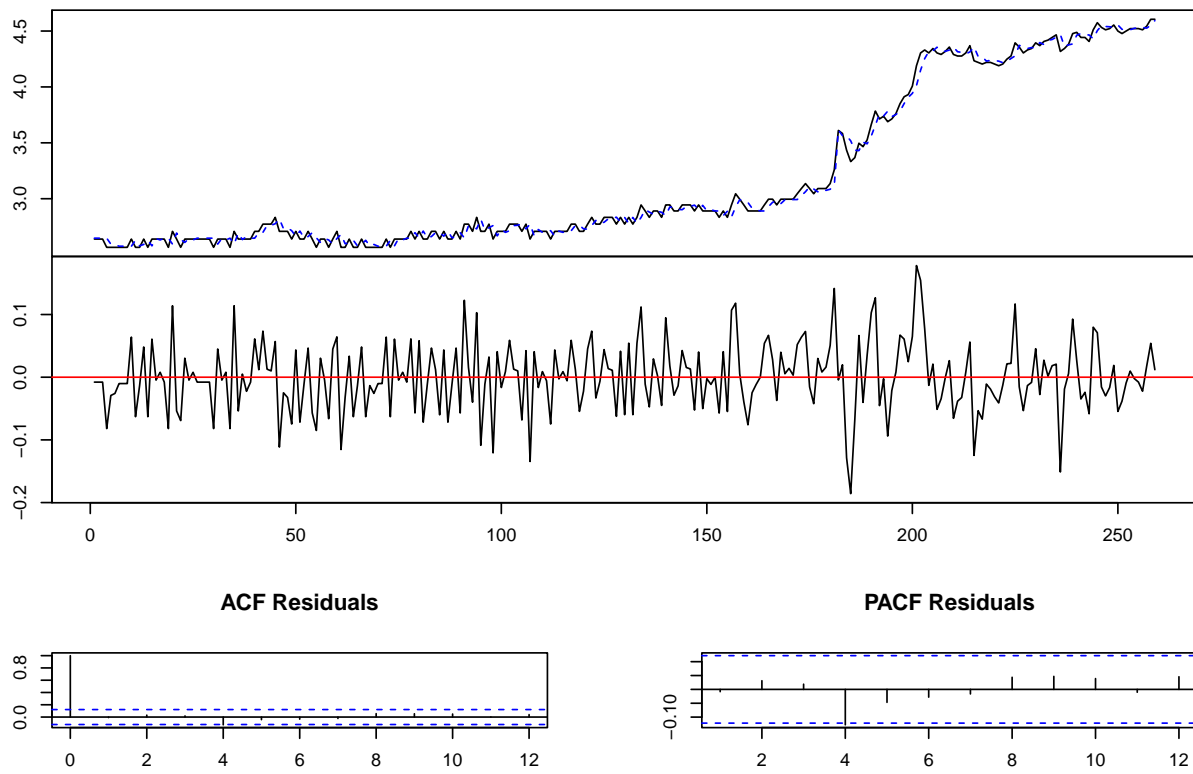
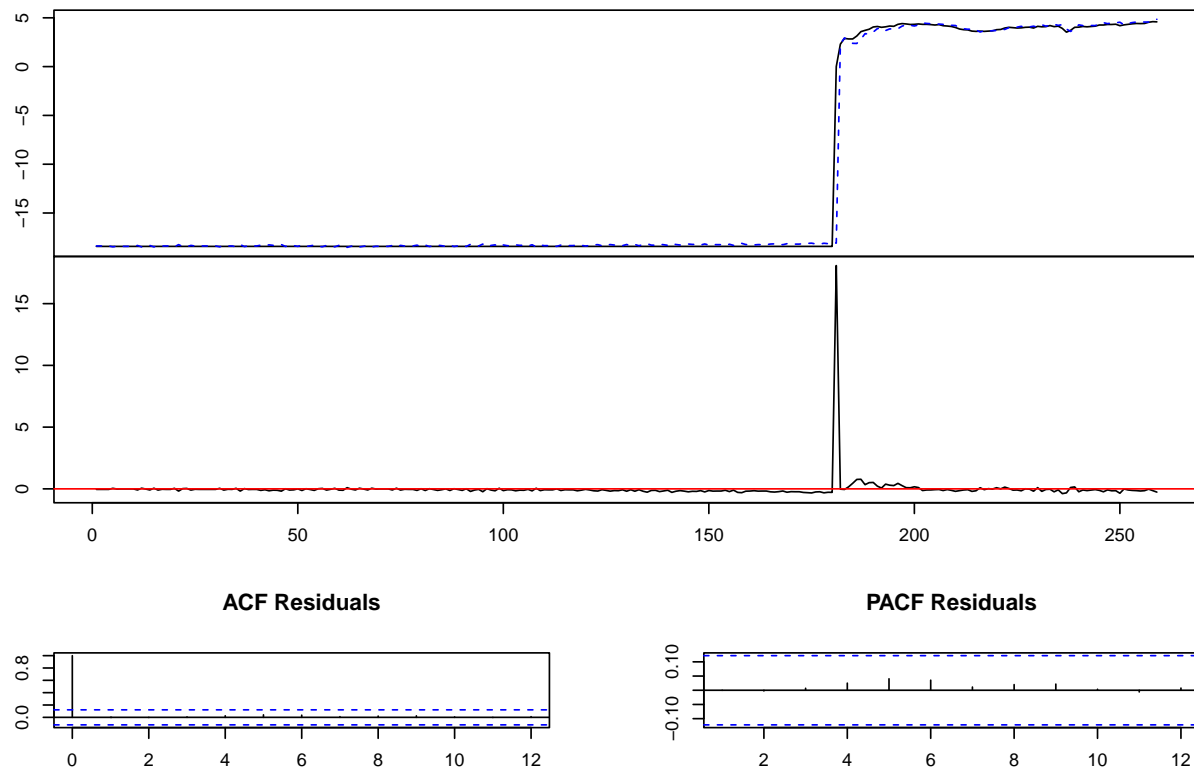


Diagram of fit and residuals for ChatGPT



Also in this case a significant spike at lag 4 of the PACF plot is detected.

Let's use the tests:

```
serial.test(series.var3)
```

Portmanteau Test (asymptotic)

```
data: Residuals of VAR object series.var3  
Chi-squared = 54.782, df = 52, p-value = 0.3695
```

```
serial.test(series.var3, type = "BG")
```


Breusch-Godfrey LM test


data: Residuals of VAR object series.var3
Chi-squared = 30.073, df = 20, p-value = 0.06868

At 5% significance level, there is no evidence of autocorrelation in the residuals, so the VAR model with 3 lags is going to be selected because of its lower BIC and also less parameters.

Impulse response functions


```
plot(irf(series.var3, n.ahead = 18))
```

Orthogonal Impulse Response from AI





95 % Bootstrap CI, 100 runs

Orthogonal Impulse Response from ChatGPT



```
plot(fevd(series.var3, n.ahead = 18))
```



Impulse response function for AI:

- **Impact on AI:** When there is a sudden spike in searches for “AI” (e.g., a significant news event related to AI), searches for “AI” itself see an immediate positive effect. This means that interest in “AI” increases sharply in response to relevant events, but the effect stabilizes quickly, indicating that the spike in interest doesn’t last long.
- **Impact on ChatGPT:** A sudden spike in searches for “AI” gradually influences searches for “ChatGPT”. The impact starts small and grows over time, suggesting that as people search for “AI”, they increasingly also search for “ChatGPT”. This could indicate that interest in AI-related developments or news gradually leads to more curiosity and searches about specific AI applications like ChatGPT.

Impulse response function from ChatGPT:

- **Impact on ChatGPT:** When there is a sudden spike in searches for “ChatGPT” (e.g., a significant news event or release related to ChatGPT), searches for “ChatGPT” itself

increase sharply but this effect gradually declines. This suggests that interest in “ChatGPT” spikes in response to specific events but fades over time as the initial excitement wears off.

- **Impact on AI:** A sudden spike in searches for “ChatGPT” has a negligible effect on searches for “AI”. This indicates that increased interest in ChatGPT does not lead to a significant increase in interest in the broader topic of AI. People searching for ChatGPT may already be aware of AI or may not feel the need to search for AI specifically.

Interpretation of FEVD:

FEVD for AI:

- The forecast error variance decomposition for “AI” shows that most of the variability in searches for “AI” can be explained by its own past search behavior. Over time, a small portion of the variability can be explained by searches for “ChatGPT”, indicating that interest in “AI” is primarily self-driven with some influence from the increasing interest in “ChatGPT”.

FEVD for ChatGPT:

- The forecast error variance decomposition for “ChatGPT” indicates that almost all of the variability in searches for “ChatGPT” is explained by its own past search behavior, with very little influence from searches for “AI”. This suggests that interest in “ChatGPT” is largely self-sustained and independent of the broader interest in AI.

VAR FORECASTING

Let’s now evaluate the forecasting power of the model:

```
series.var3.forecast <- predict(series.var3,
                                n.ahead = 14,
                                ci = 0.95) # 95% confidence interval
```

Results are stored in a data frame:

```
gpt_var_forecast <- xts(series.var3.forecast$fcst$ChatGPT[,-4],
                        # we exclude the last column with CI
                        tail(index(series), 14))
ai_var_forecast <- xts(series.var3.forecast$fcst$AI[,-4],
                      # we exclude the last column with CI
                      tail(index(series), 14))
```

```
names(gpt_var_forecast) <- c("gpt_fore", "gpt_lower", "gpt_upper")
names(ai_var_forecast) <- c("ai_fore", "ai_lower", "ai_upper")
```


```
ai_gpt_forecasts <- merge(series,
                           gpt_var_forecast,
                           ai_var_forecast)
```

Let's now inspect them visually:


```
plot(ai_gpt_forecasts["2022-12-04/", c("AI", "ai_fore",
                                         "ai_lower", "ai_upper")],
     major.ticks = "months",
     grid.ticks.on = "months",
     grid.ticks.lty = 7,
     main = "14 weeks forecast of AI",
     col = c("black", "blue", "red", "red"))
```

14 weeks forecast of AI

2022-12-04 / 2024-05-26



```
plot(ai_gpt_forecasts["2022-12-04/", c("ChatGPT", "gpt_fore")],  
     major.ticks = "months",  
     grid.ticks.on = "months",  
     grid.ticks.lty = 7,  
     main = "14 weeks forecast of ChatGPT",  
     col = c("black", "blue"))
```



From the plots it is already pretty clear that the VAR model is performing worse than the ARIMA alternatives, but let's compare the error metrics anyway.

```
ai_gpt_forecasts$AI <- exp(ai_gpt_forecasts$AI)
ai_gpt_forecasts$ai_fore <- exp(ai_gpt_forecasts$ai_fore)
ai_gpt_forecasts$ai_lower <- exp(ai_gpt_forecasts$ai_lower)
ai_gpt_forecasts$ai_upper <- exp(ai_gpt_forecasts$ai_upper)
ai_gpt_forecasts$mae.ai <- abs(ai_gpt_forecasts$AI -
  ↪ ai_gpt_forecasts$ai_fore)
ai_gpt_forecasts$mse.ai <- (ai_gpt_forecasts$AI -
  ↪ ai_gpt_forecasts$ai_fore) ^ 2
ai_gpt_forecasts$mape.ai <- abs((ai_gpt_forecasts$AI -
  ↪ ai_gpt_forecasts$ai_fore)/ai_gpt_forecasts$AI)
```

```

ai_gpt_forecasts$amape.ai <- abs((ai_gpt_forecasts$AI -
  ↪ ai_gpt_forecasts$ai_fore) / (ai_gpt_forecasts$AI +
  ↪ ai_gpt_forecasts$ai_fore))

ai_gpt_forecasts$ChatGPT <- exp(ai_gpt_forecasts$ChatGPT)
ai_gpt_forecasts$gpt_fore <- exp(ai_gpt_forecasts$gpt_fore)
ai_gpt_forecasts$gpt_lower <- exp(ai_gpt_forecasts$gpt_lower)
ai_gpt_forecasts$gpt_upper <- exp(ai_gpt_forecasts$gpt_upper)
ai_gpt_forecasts$mae.gpt <- abs(ai_gpt_forecasts$ChatGPT -
  ↪ ai_gpt_forecasts$gpt_fore)
ai_gpt_forecasts$mse.gpt <- (ai_gpt_forecasts$ChatGPT -
  ↪ ai_gpt_forecasts$gpt_fore) ^ 2
ai_gpt_forecasts$mape.gpt <- abs((ai_gpt_forecasts$ChatGPT -
  ↪ ai_gpt_forecasts$gpt_fore)/ai_gpt_forecasts$ChatGPT)
ai_gpt_forecasts$amape.gpt <- abs((ai_gpt_forecasts$ChatGPT -
  ↪ ai_gpt_forecasts$gpt_fore) / (ai_gpt_forecasts$ChatGPT +
  ↪ ai_gpt_forecasts$gpt_fore))

var_metrics = colMeans(ai_gpt_forecasts[,9:16], na.rm = TRUE)

```

FINAL COMPARISON

Let's now compare the models' forecasting power, by starting with **ChatGPT** time series:

```

ts2_combined_perf$VAR <- (var_metrics[5:8])
ts2_combined_perf

```

	ARIMA_1_1_1	ARIMA_0_1_1	VAR
mae	9.59507686	9.85323646	5.962454e+02
mse	167.26598112	174.59323412	6.274502e+05
mape	0.10908605	0.11202213	6.799866e+00
amape	0.05960564	0.06137572	6.660865e-01

- The table clearly shows that ARIMA(1,1,1) is the best at forecasting future values compared to the others.
- VAR model has substantially higher errors across all metrics, suggesting that it is not well-suited for the given data.

Let's also examine **AI** time series performances table:

```
combined_perf$VAR <- (var_metrics[1:4])
combined_perf
```

	ARIMA_1_1_2	ARIMA_0_1_8	VAR
mae	3.65165914	4.53356390	14.43073561
mse	19.48744607	31.70184775	224.24582515
mape	0.04219822	0.05237733	0.15589074
amape	0.02062693	0.02519301	0.07190771

- The table shows that ARIMA(1,1,2) is the best model at forecasting future values compared to the others.
- Also in this case, VAR model has higher errors across all metrics, suggesting that it is not well-suited for the given data.

Justification and Motivation for Results

1. ARIMA Model Strengths:

- **Univariate Nature:** ARIMA models are designed for univariate time series data, making them ideal when the primary goal is to predict the future values of a single series without considering cross-series interactions.
- **Strong Performance on Trend Data:** Google Trends data often exhibit clear patterns and trends, which ARIMA models can effectively capture and extrapolate.

2. VAR Model Weaknesses:

- **Complexity and Data Requirements:** VAR models are more complex and require more data to accurately estimate the relationships between multiple time series. If these relationships are weak or not properly captured, the model's performance can degrade significantly.
- **Overfitting Risks:** With more parameters to estimate, VAR models are prone to overfitting, especially if the underlying data does not strongly support the assumed interdependencies.

Final conclusions

This project demonstrates the importance of model selection and validation in time series forecasting. While ARIMA models have shown strong performance with Google Trends data, the use of VAR models requires careful consideration of data characteristics and model specifications. The results obtained justify prioritizing ARIMA models for similar forecasting tasks, ensuring reliable and accurate predictions.